

# Depth-Supervised NeRF: Fewer Views and Faster Training for Free

(CVPR 2022)

한국과학기술연구원(KIST)

CVIPL 학생연구원 김연욱



한국과학기술연구원  
Korea Institute of Science and Technology

# Depth-supervised NeRF: Fewer Views and Faster Training for Free

Kangle Deng<sup>1</sup>

Andrew Liu<sup>2</sup>

Jun-Yan Zhu<sup>1</sup>

Deva Ramanan<sup>1,3</sup>

<sup>1</sup>Carnegie Mellon University

<sup>2</sup>Google

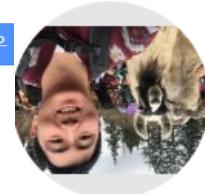
<sup>3</sup>Argo AI



Kangle Deng

[Carnegie Mellon University](#)  
andrew.cmu.edu의 이메일 확인됨 - [홈페이지](#)  
Computer Vision

팔로우



Andrew Liu

[Google](#)  
google.com의 이메일 확인됨 - [홈페이지](#)  
3D Computer Vision Generative Models

팔로우

제목	인용	연도	제목	인용	연도
Depth-supervised nerf: Fewer views and faster training for free K Deng, A Liu, JY Zhu, D Ramanan Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern ...	1131	2022	Depth-supervised nerf: Fewer views and faster training for free K Deng, A Liu, JY Zhu, D Ramanan Proceedings of the IEEE/CVF conference on computer vision and pattern ...	1130	2022
IRC-GAN: Introspective Recurrent Convolutional GAN for Text-to-video Generation. K Deng, T Fei, X Huang, Y Peng International Joint Conferences on Artificial Intelligence (IJCAI) 2019	61	2019	Fighting fake news: Image splice detection via learned self-consistency M Huh, A Liu, A Owens, AA Efros Proceedings of the European conference on computer vision (ECCV), 101-117	573	2018
3d-aware conditional image synthesis K Deng, G Yang, D Ramanan, JY Zhu Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern ...	50	2023	Urban radiance fields K Rematas, A Liu, PP Srinivasan, JT Barron, A Tagliasacchi, ... Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern ...	383	2022

# Depth-supervised NeRF: Fewer Views and Faster Training for Free

Kangle Deng<sup>1</sup>

<sup>1</sup>Carnegie Mellon University

Andrew Liu<sup>2</sup>

Jun-Yan Zhu<sup>1</sup>

<sup>2</sup>Google

Deva Ramanan<sup>1,3</sup>

<sup>3</sup>Argo AI



Jun-Yan Zhu

Assistant Professor, [Carnegie Mellon University](#)  
cs.cmu.edu의 이메일 확인됨 - [홈페이지](#)

Computer Vision Computer Graphics Generative Models Computational Photography

팔로우



Deva Ramanan

Professor, Robotics Institute, [Carnegie Mellon University](#).  
cs.cmu.edu의 이메일 확인됨 - [홈페이지](#)  
Computer Vision Machine Learning

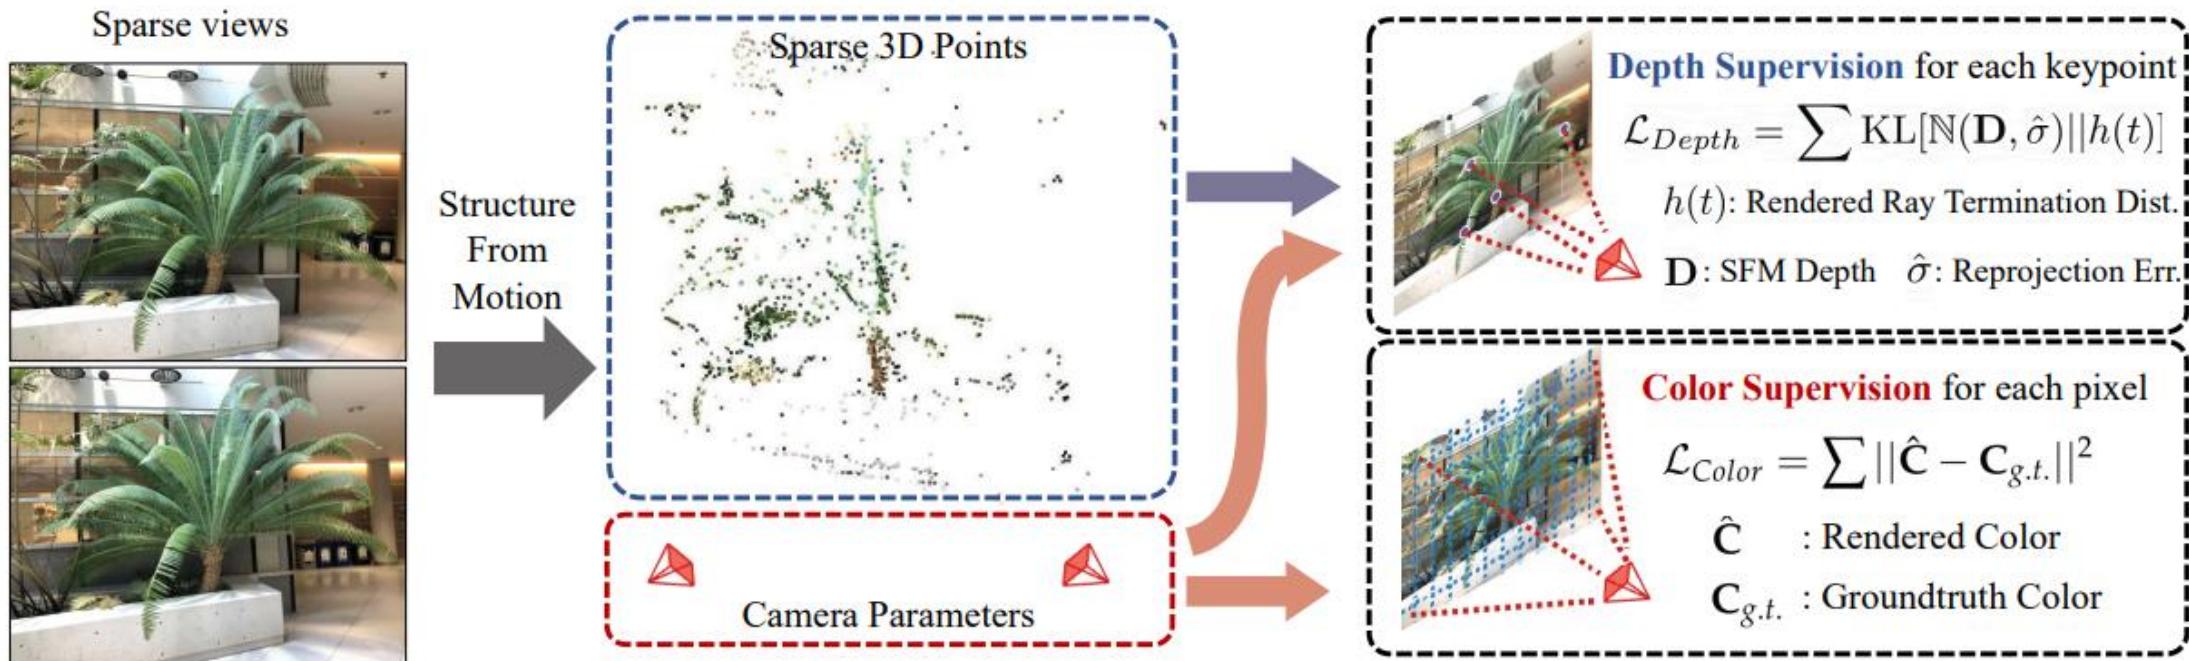
팔로우

제목	인용	연도
<a href="#">Unpaired image-to-image translation using cycle-consistent adversarial networks</a> JY Zhu, T Park, P Isola, AA Efros Proceedings of the IEEE International Conference on Computer Vision	29090	2017
<a href="#">Image-to-image translation with conditional adversarial networks</a> P Isola, JY Zhu, T Zhou, AA Efros Proceedings of the IEEE Conference on Computer Vision and Pattern ...	28408	2017
<a href="#">High-Resolution Image Synthesis and Semantic Manipulation with Conditional GANs</a> TC Wang, MY Liu, JY Zhu, A Tao, J Kautz, B Catanzaro Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition	5669	2018

제목	인용	연도
<a href="#">Microsoft coco: Common objects in context</a> TY Lin, M Maire, S Belongie, J Hays, P Perona, D Ramanan, P Dollár, ... European conference on computer vision, 740-755	62768	2014
<a href="#">Object detection with discriminatively trained part-based models</a> PF Felzenszwalb, RB Girshick, D McAllester, D Ramanan IEEE transactions on pattern analysis and machine intelligence 32 (9), 1627-1645	13612	2009

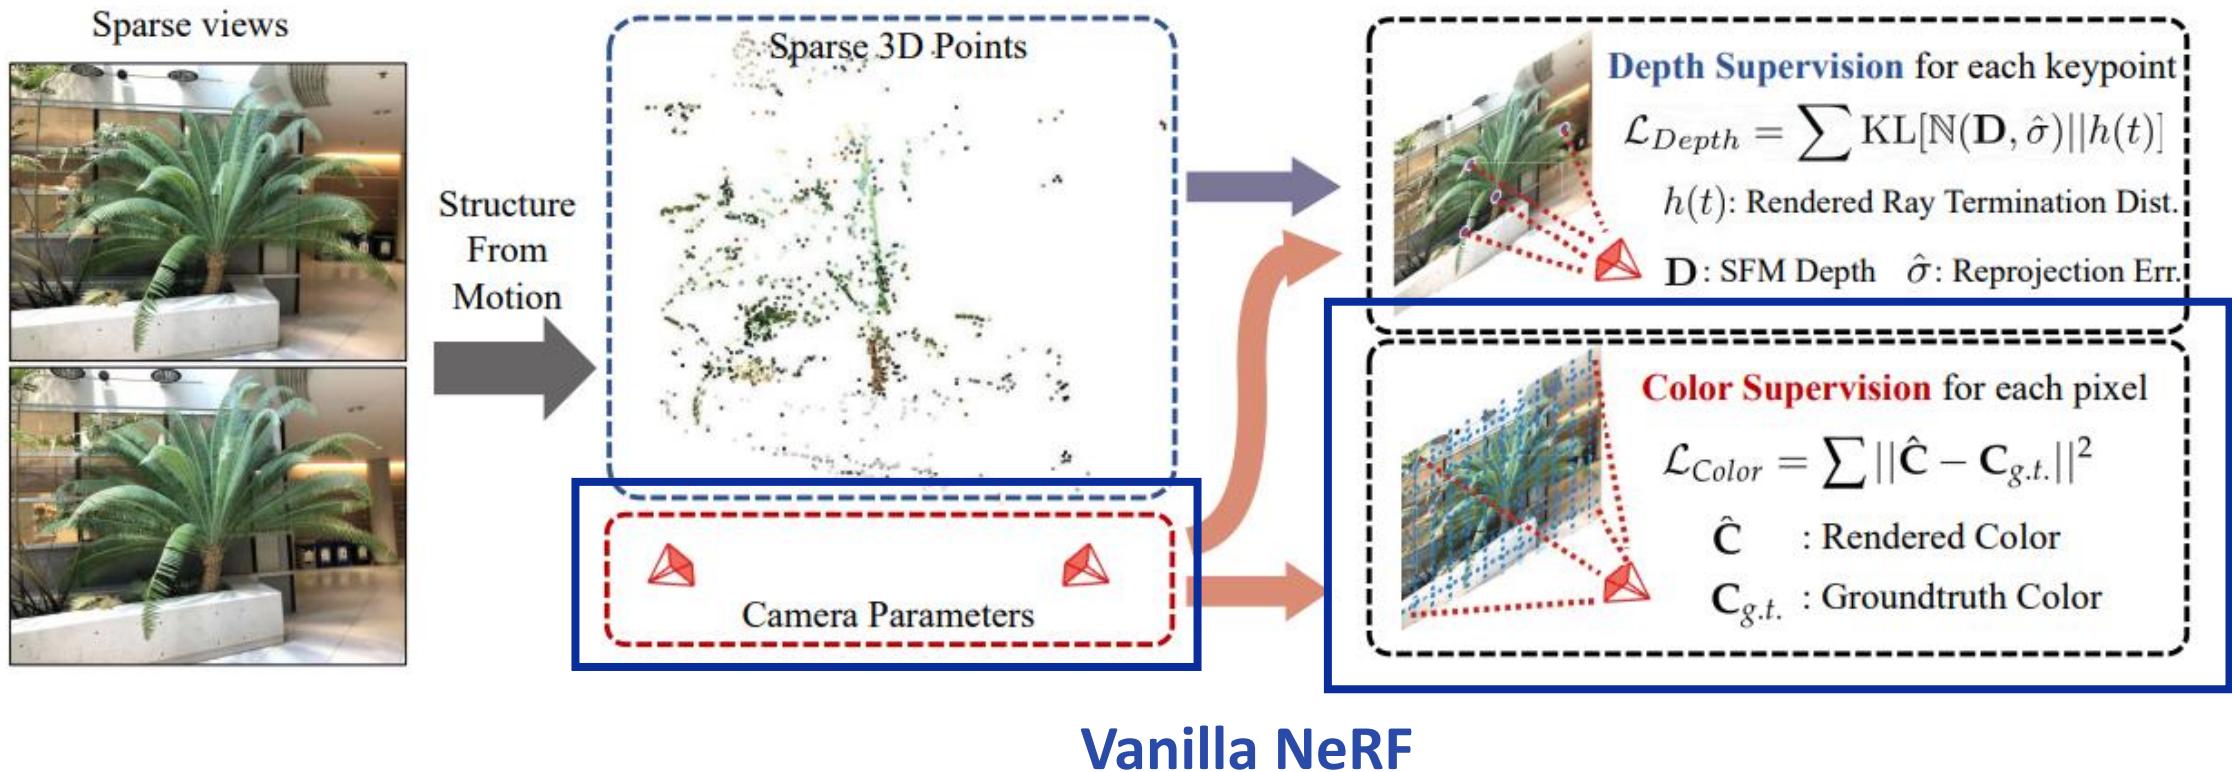
# Depth-supervised NeRF

## Overview



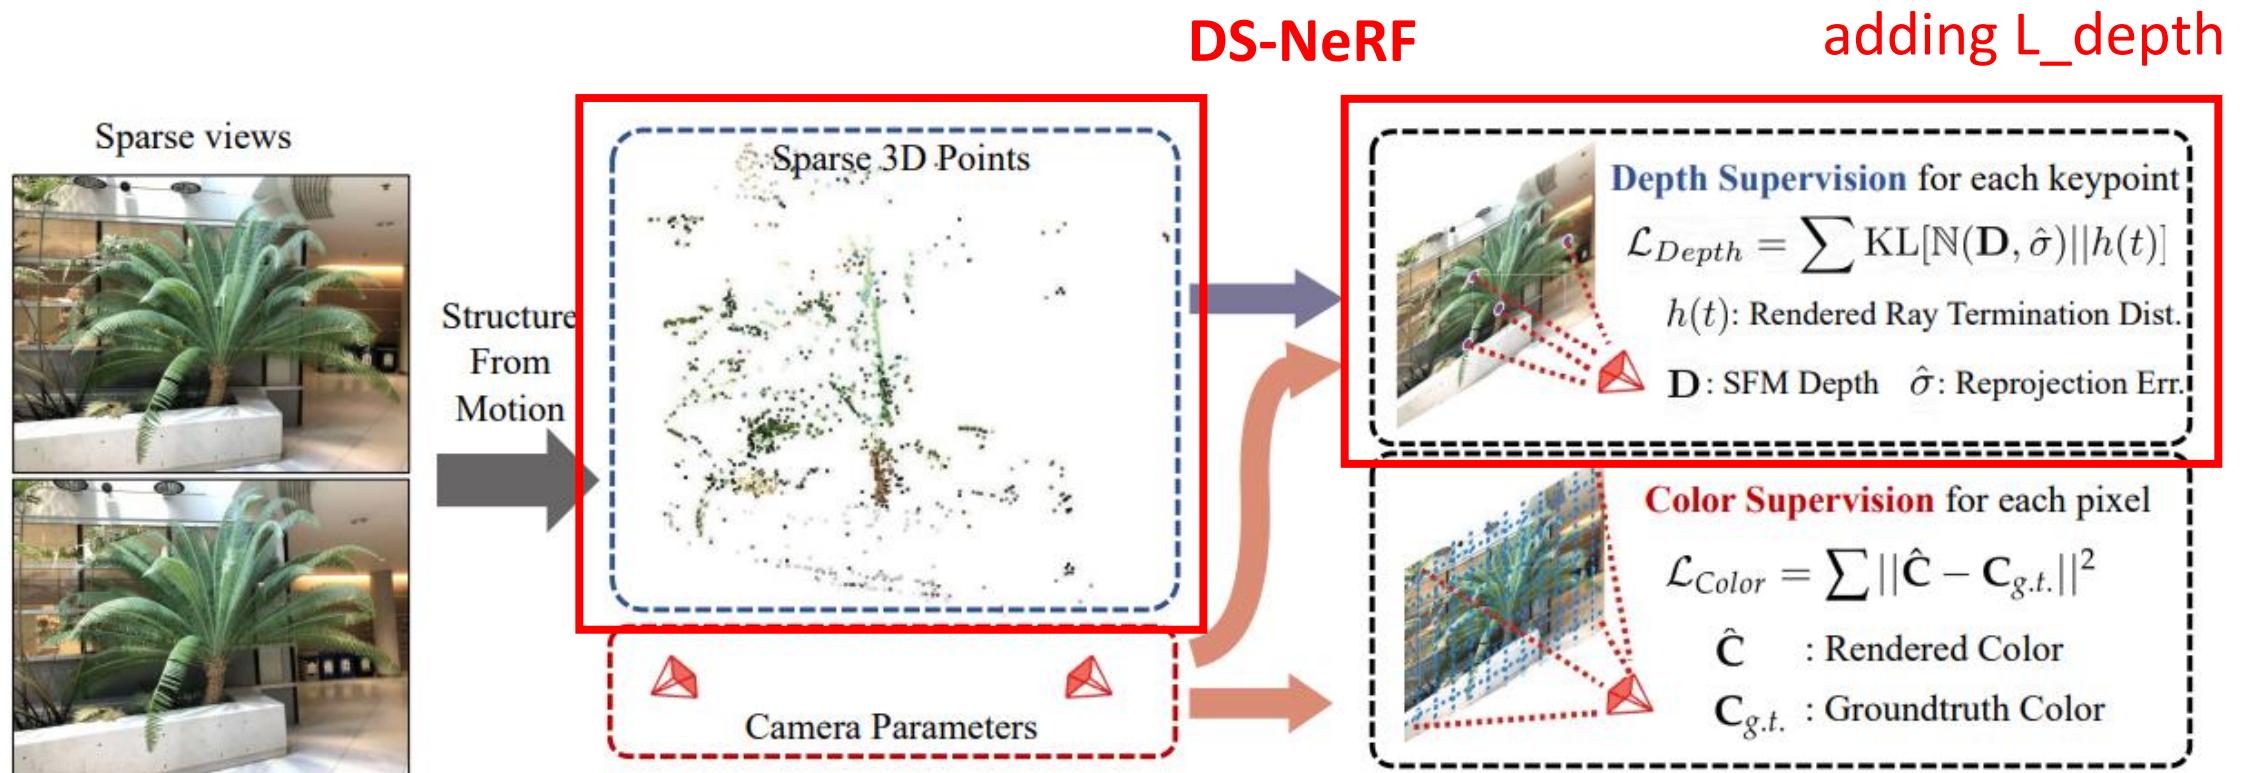
# Depth-supervised NeRF

## Overview



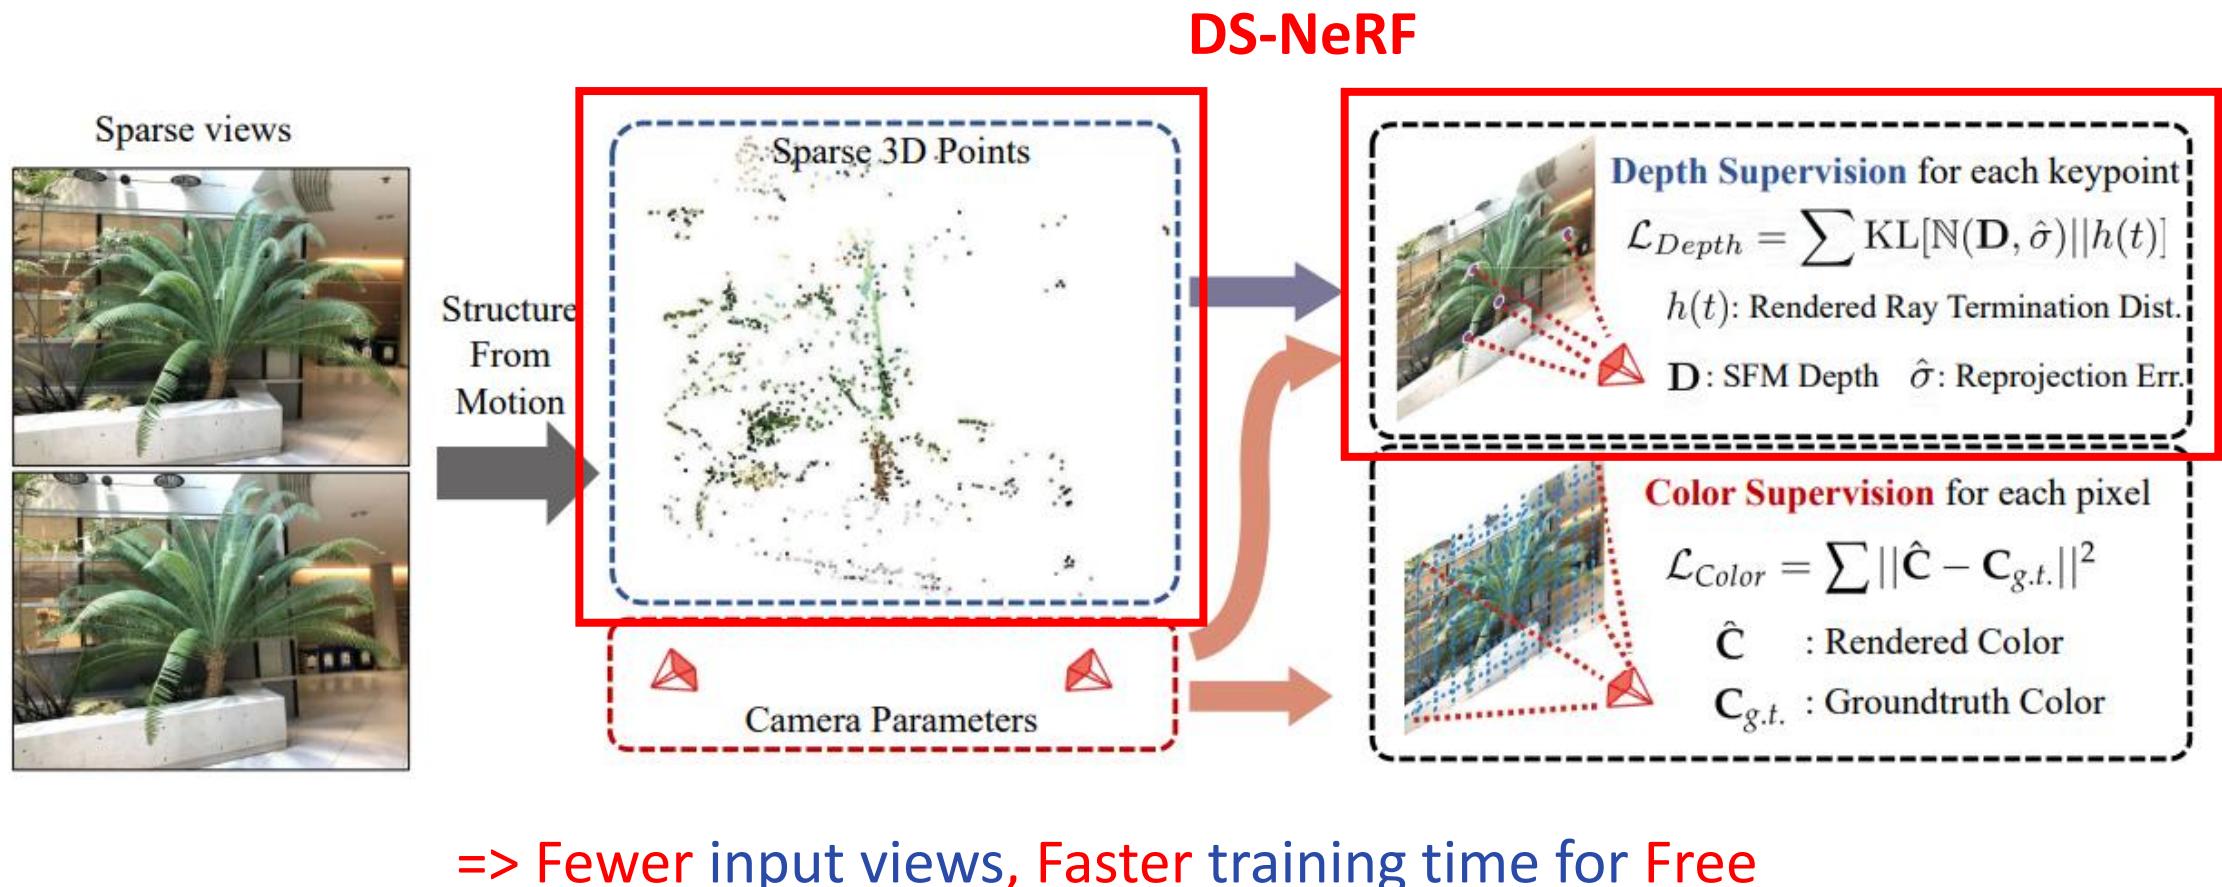
# Depth-supervised NeRF

## Overview



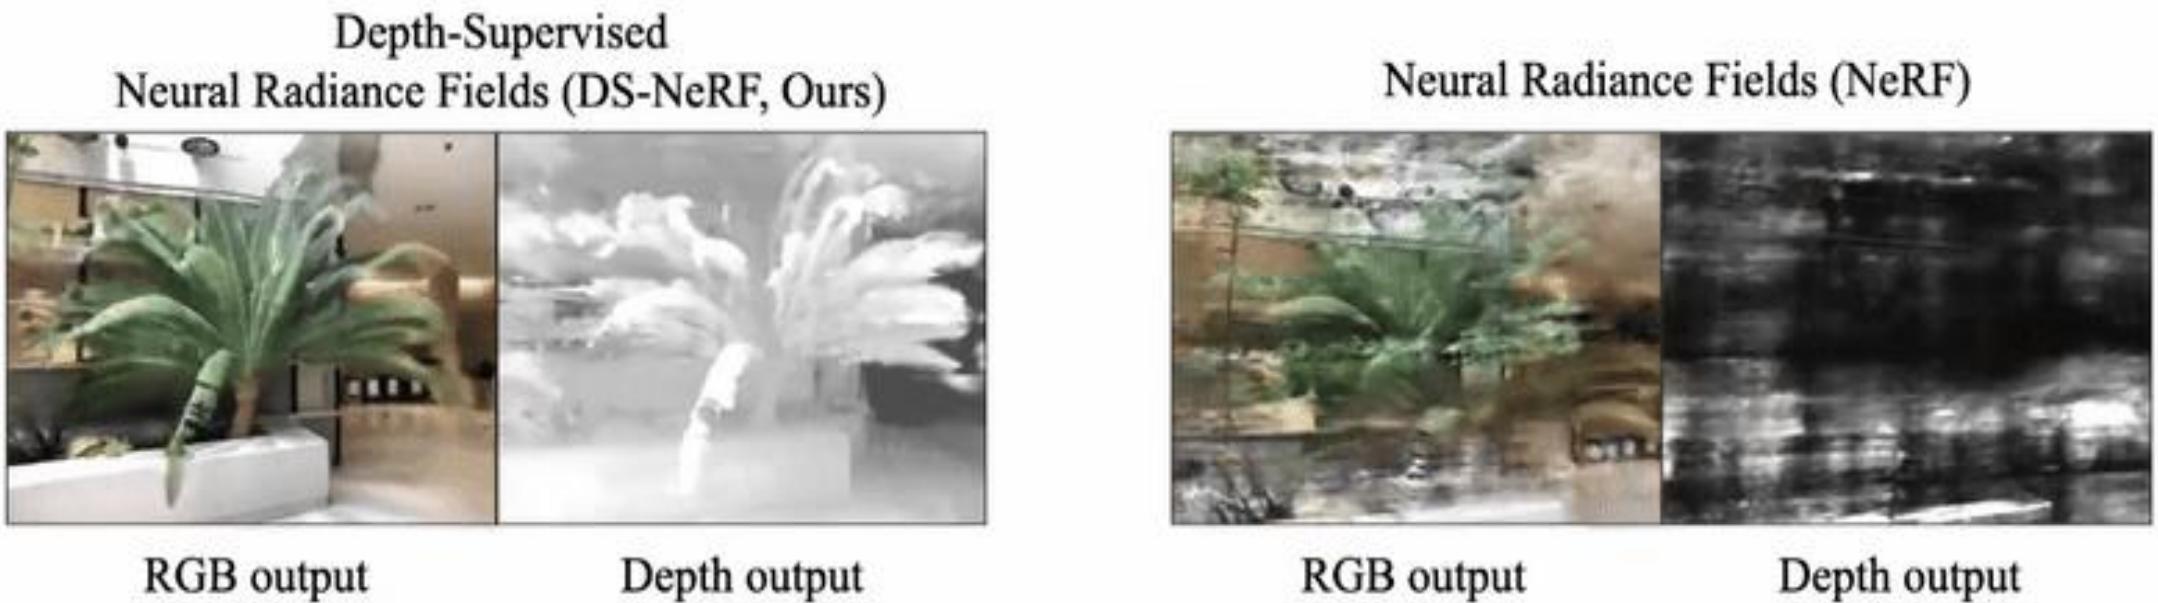
# Depth-supervised NeRF

## Overview



## Depth-supervised NeRF

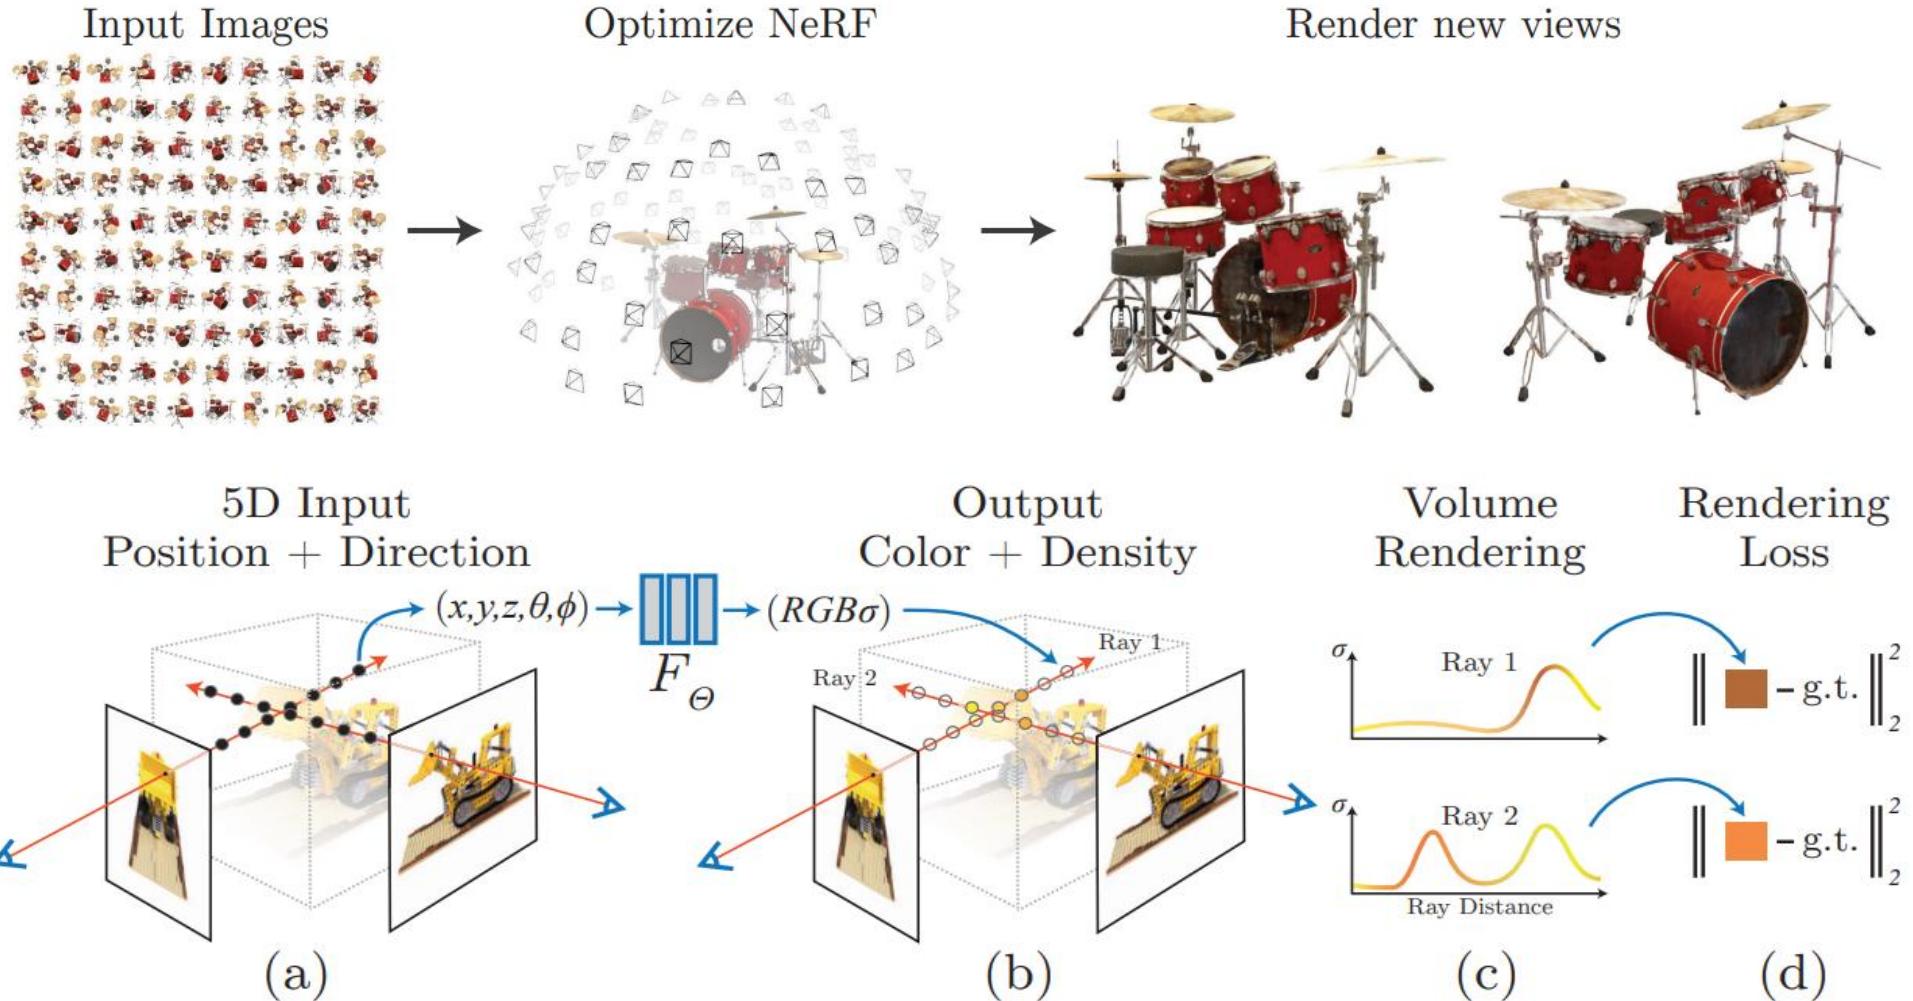
---



DS-NeRF vs NeRF when two input views

# Motivation

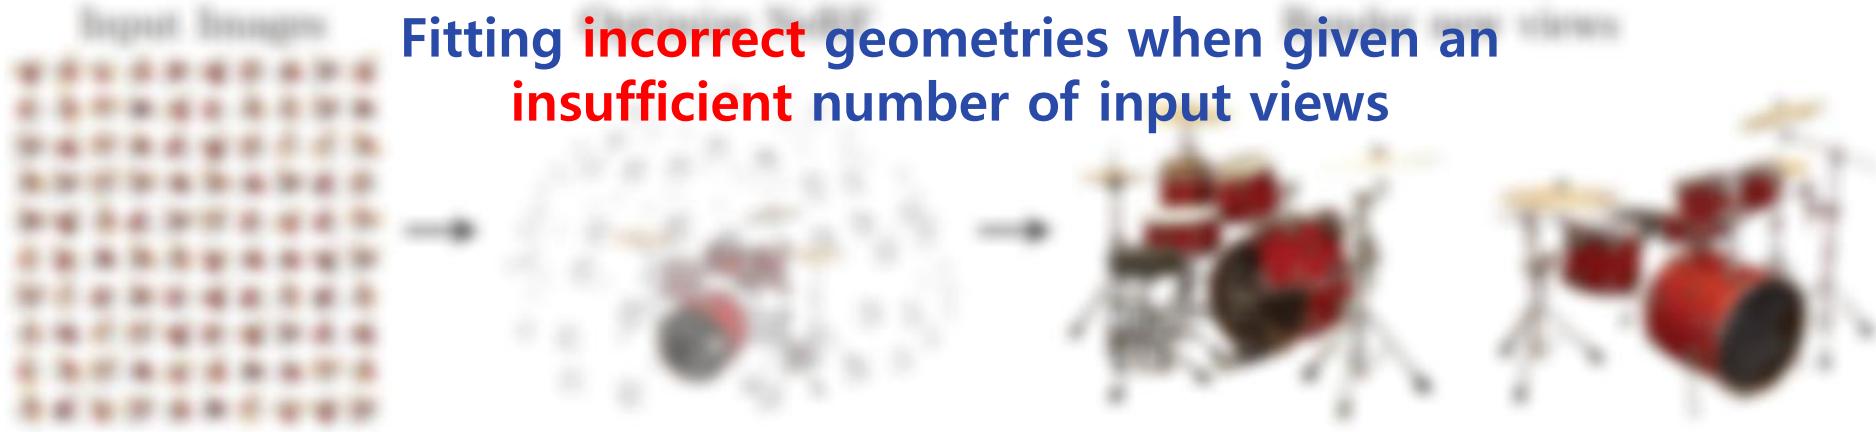
# Vanilla NeRF's Problem



## Vanilla NeRF's Problem

---

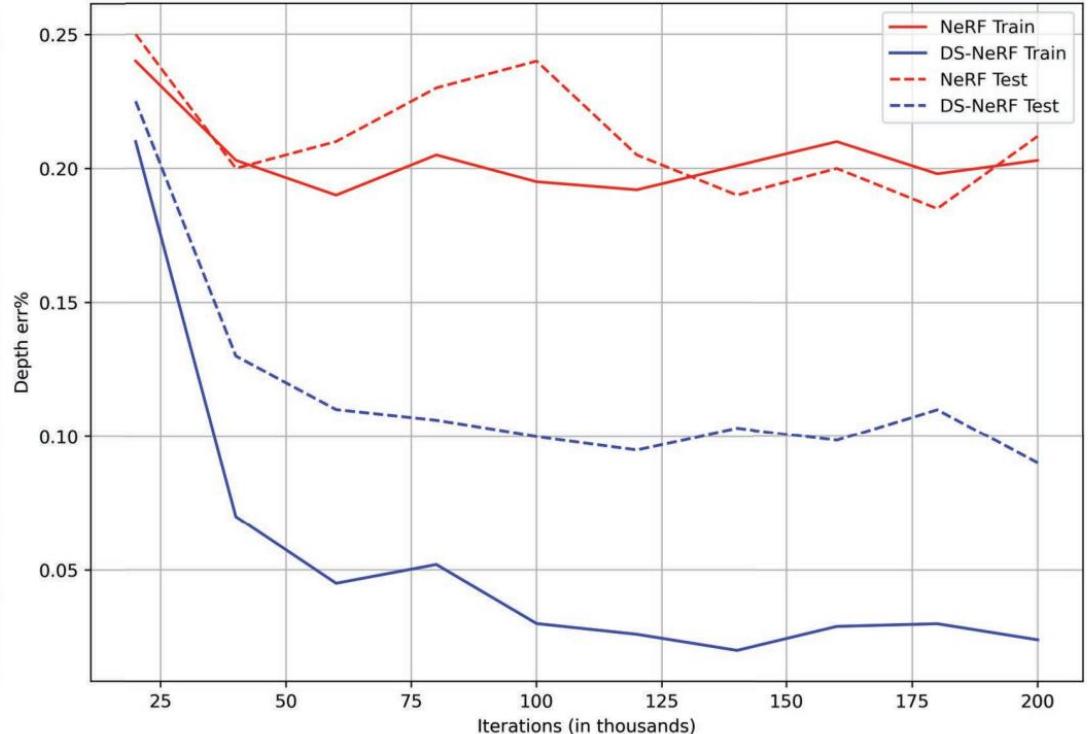
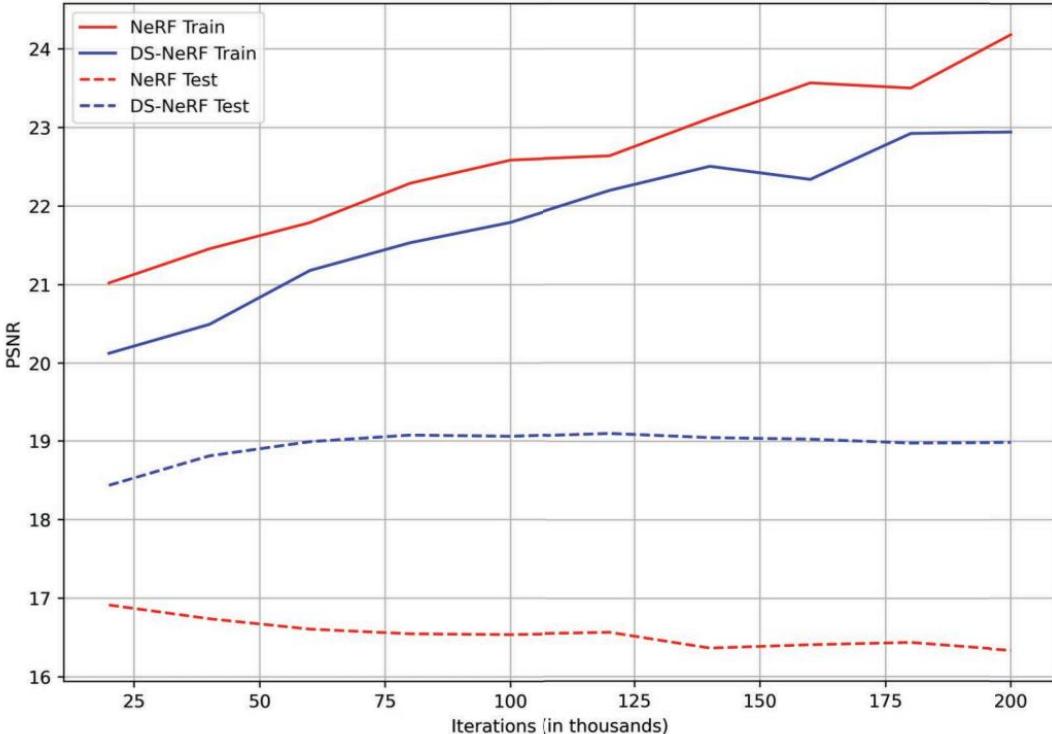
Fitting **incorrect** geometries when given an  
**insufficient** number of input views



Standard volumetric rendering does **not enforce**  
the constraint that most of a scene's geometry  
consist of **empty space** and **opaque surfaces**.  
=> Training Time ↑

## Vanilla NeRF's Problem

---



NeRF is susceptible to **overfitting** when given few training views.

# Structure From Motion (COLMAP)

---

## COLMAP Output

### cameras.txt %

This file contains the intrinsic parameters of all reconstructed cameras in the dataset using one line per camera, e.g.:

```
# Camera List with one line of data per camera:  
#   CAMERA_ID, MODEL, WIDTH, HEIGHT, PARAMS[]  
# Number of cameras: 3  
1 SIMPLE_PINHOLE 3072 2304 2559.81 1536 1152  
2 PINHOLE 3072 2304 2560.56 2560.56 1536 1152  
3 SIMPLE_RADIAL 3072 2304 2559.69 1536 1152 -0.0218531
```

### points3D.txt

This file contains the information of all reconstructed 3D points in the dataset using one line per point, e.g.:

```
# 3D point List with one line of data per point:  
#   POINT3D_ID, X, Y, Z, R, G, B, ERROR, TRACK[] as (IMAGE_ID, POINT2D_IDX)  
# Number of points: 3, mean track Length: 3.3334  
63390 1.67241 0.292931 0.609726 115 121 122 1.33927 16 6542 15 7345 6 6714 14 7227  
63376 2.01848 0.108877 -0.0260841 102 209 250 1.73449 16 6519 15 7322 14 7212 8 3991  
63371 1.71102 0.28566 0.53475 245 251 249 0.612829 118 4140 117 4473
```

### images.txt %

This file contains the pose and keypoints of all reconstructed images in the dataset using two lines per image, e.g.:

```
# Image List with two Lines of data per image:  
#   IMAGE_ID, QW, QX, QY, QZ, TX, TY, TZ, CAMERA_ID, NAME  
#   POINTS2D[] as (X, Y, POINT3D_ID)  
# Number of images: 2, mean observations per image: 2  
1 0.851773 0.0165051 0.503764 -0.142941 -0.737434 1.02973 3.74354 1 P1180141.JPG  
2362.39 248.498 58396 1784.7 268.254 59027 1784.7 268.254 -1  
2 0.851773 0.0165051 0.503764 -0.142941 -0.737434 1.02973 3.74354 1 P1180142.JPG  
1190.83 663.957 23056 1258.77 640.354 59070
```

### rigs.txt

This file contains the configured rigs and sensors, e.g.:

```
# Rig calib list with one line of data per calib:  
#   RIG_ID, NUM_SENSORS, REF_SENSOR_TYPE, REF_SENSOR_ID, SENSORS[] as (SENSOR_TYPE, SENSOR_ID, HAS_POSE,  
#   Number of rigs: 1  
1 2 CAMERA 1 CAMERA 2 1 -0.9999701516465348 -0.0011120266840749639 -0.0075347911527510894 0.001298512589  
2 1 CAMERA 3
```

# Structure From Motion (COLMAP)

## COLMAP Output

### cameras.txt

This file contains the intrinsic parameters of all reconstructed cameras in the dataset using one line per camera, e.g.:

```
# Camera list with one line of data per camera:  
#   CAMERA_ID, MODEL, WIDTH, HEIGHT, PARAMS[]  
# Number of cameras: 3  
1 SIMPLE_PINHOLE 3072 2304 2559.81 1536 1152  
2 PINHOLE 3072 2304 2560.56 2560.56 1536 1152  
3 SIMPLE_RADIAL 3072 2304 2559.69 1536 1152 -0.0218531
```

### images.txt

This file contains the pose and keypoints of all reconstructed images in the dataset using two lines per image, e.g.:

```
# Image list with two lines of data per image:  
#   IMAGE_ID, QW, QX, QY, QZ, TX, TY, TZ, CAMERA_ID, NAME  
#   POINTS2D[] as (X, Y, POINT3D_ID)  
# Number of images: 2, mean observations per image: 2  
1 0.851773 0.0165051 0.503764 -0.142941 -0.737434 1.02973 3.74354 1 P1180141.JPG  
2362.39 248.498 58396 1784.7 268.254 59027 1784.7 268.254 -1  
2 0.851773 0.0165051 0.503764 -0.142941 -0.737434 1.02973 3.74354 1 P1180142.JPG  
1190.83 663.957 23056 1258.77 640.354 59070
```

### points3D.txt

This file contains the information of all reconstructed 3D points in the dataset using one line per point, e.g.:

```
# 3D point list with one line of data per point:  
#   POINT3D_ID, X, Y, Z, R, G, B, ERROR, TRACK[] as (IMAGE_ID, POINT2D_IDX)  
# Number of points: 3, mean track Length: 3.3334  
63390 1.67241 0.292931 0.609726 115 121 122 1.33927 16 6542 15 7345 6 6714 14 7227  
63376 2.01848 0.108877 -0.0260841 102 209 250 1.73449 16 6519 15 7322 14 7212 8 3991  
63371 1.71102 0.28566 0.53475 245 251 249 0.612829 118 4140 117 4473
```

### rigs.txt

This file contains the configured rigs and sensors, e.g.:

```
# Rig calib list with one line of data per calib:  
#   RIG_ID, NUM_SENSORS, REF_SENSOR_TYPE, REF_SENSOR_ID, SENSORS[] as (SENSOR_TYPE, SENSOR_ID, HAS_POSE,  
#   Number of rigs: 1  
1 2 CAMERA 1 CAMERA 2 1 -0.9999701516465348 -0.0011120266840749639 -0.0075347911527510894 0.001298512589  
2 1 CAMERA 3
```

# Structure From Motion (COLMAP)

## COLMAP Output

### cameras.txt

This file contains the intrinsic parameters of all reconstructed cameras in the dataset using one line per camera, e.g.:

```
# Camera list with one line of data per camera:  
#   CAMERA_ID, MODEL, WIDTH, HEIGHT, PARAMS[]  
# Number of cameras: 3  
1 SIMPLE_PINHOLE 3072 2304 2559.81 1536 1152  
2 PINHOLE 3072 2304 2560.56 2560.56 1536 1152  
3 SIMPLE_RADIAL 3072 2304 2559.69 1536 1152 -0.0218531
```

### images.txt

This file contains the pose and keypoints of all reconstructed images in the dataset using two lines per image, e.g.:

```
# Image list with two lines of data per image:  
#   IMAGE_ID, QW, QX, QY, QZ, TX, TY, TZ, CAMERA_ID, NAME  
#   POINTS2D[] as (X, Y, POINT3D_ID)  
# Number of images: 2, mean observations per image: 2  
1 0.851773 0.0165051 0.503764 -0.142941 -0.737434 1.02973 3.74354 1 P1180141.JPG  
2362.39 248.498 58396 1784.7 268.254 59027 1784.7 268.254 -1  
2 0.851773 0.0165051 0.503764 -0.142941 -0.737434 1.02973 3.74354 1 P1180142.JPG  
1190.83 663.957 23056 1258.77 640.354 59070
```

## Sparse 3D Point Cloud(의 Z값 = Depth)

### points3D.txt

This file contains the information of all reconstructed 3D points in the dataset using one line per point, e.g.:

재투영오차

```
# 3D point list with one line of data per point:  
#   POINT3D_ID, X, Y, Z, R, G, B, ERROR, TRACK[] as (IMAGE_ID, POINT2D_IDX)  
# Number of points: 3, mean track length: 3.3334  
63390 1.67241 0.292931 0.609726 115 121 122 1.33927 16 6542 15 7345 6 6714 14 7227  
63376 2.01848 0.108877 -0.0260841 102 209 250 1.73449 16 6519 15 7322 14 7212 8 3991  
63371 1.71102 0.28566 0.53475 245 251 249 0.612829 118 4140 117 4473
```

### rigs.txt

This file contains the configured rigs and sensors, e.g.:

```
# Rig calib list with one line of data per calib:  
#   RIG_ID, NUM_SENSORS, REF_SENSOR_TYPE, REF_SENSOR_ID, SENSORS[] as (SENSOR_TYPE, SENSOR_ID, HAS_POSE,  
#   Number of rigs: 1  
1 2 CAMERA 1 CAMERA 2 1 -0.9999701516465348 -0.0011120266840749639 -0.0075347911527510894 0.001298512589  
2 1 CAMERA 3
```

# Structure From Motion (COLMAP)

## COLMAP Output

### cameras.txt %

This file contains the intrinsic parameters of all reconstructed cameras in the dataset using one line per camera, e.g.:

```
# Camera List with one line of data per camera:  
#   CAMERA_ID, MODEL, WIDTH, HEIGHT, PARAMS[]  
# Number of cameras: 3  
1 SIMPLE_PINHOLE 3072 2304 2559.81 1536 1152  
2 PINHOLE 3072 2304 2560.56 2560.56 1536 1152  
3 SIMPLE_RADIAL 3072 2304 2559.69 1536 1152 -0.0218531
```

### images.txt %

This file contains the pose and keypoints of all reconstructed images in the dataset using two lines per image, e.g.:

```
# Image List with two Lines of data per image:  
#   IMAGE_ID, QW, QX, QY, QZ, TX, TY, TZ, CAMERA_ID, NAME  
#   POINTS2D[] as (X, Y, POINT3D_ID)  
# Number of images: 2, mean observations per image: 2  
1 0.851773 0.0165051 0.503764 -0.142941 -0.737434 1.02973 3.74354 1 P1180141.JPG  
2362.39 248.498 58396 1784.7 268.254 59027 1784.7 268.254 -1  
2 0.851773 0.0165051 0.503764 -0.142941 -0.737434 1.02973 3.74354 1 P1180142.JPG  
1190.83 663.957 23056 1258.77 640.354 59070
```

for Free

### points3D.txt

This file contains the information of all reconstructed 3D points in the dataset using one line per point, e.g.:

```
# 3D point List with one line of data per point:  
#   POINT3D_ID, X, Y, Z, R, G, B, ERROR, TRACK[] as (IMAGE_ID, POINT2D_IDX)  
# Number of points: 3, mean track Length: 3.3334  
63390 1.67241 0.292931 0.609726 115 121 122 1.33927 16 6542 15 7345 6 6714 14 7227  
63376 2.01848 0.108877 -0.0260841 102 209 250 1.73449 16 6519 15 7322 14 7212 8 3991  
63371 1.71102 0.28566 0.53475 245 251 249 0.612829 118 4140 117 4473
```

### rigs.txt

This file contains the configured rigs and sensors, e.g.:

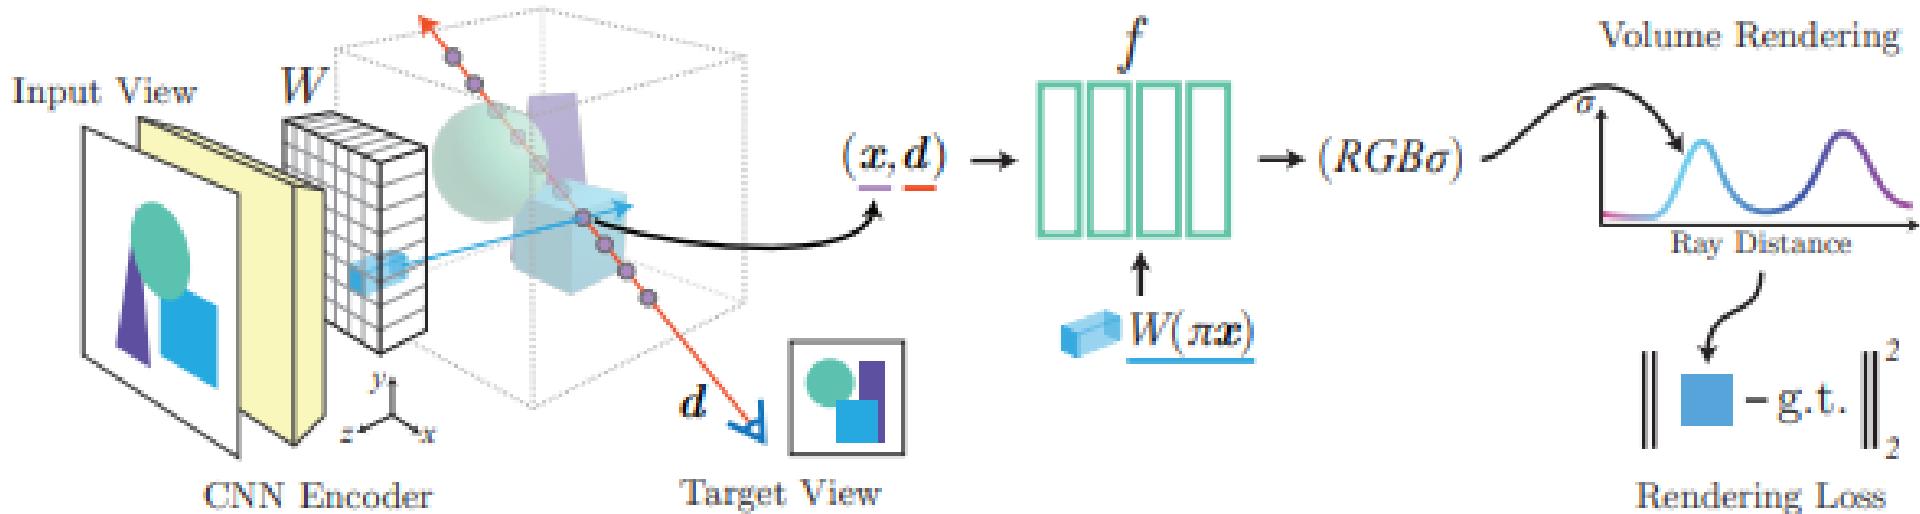
```
# Rig calib list with one line of data per calib:  
#   RIG_ID, NUM_SENSORS, REF_SENSOR_TYPE, REF_SENSOR_ID, SENSORS[] as (SENSOR_TYPE, SENSOR_ID, HAS_POSE,  
#   Number of rigs: 1  
1 2 CAMERA 1 CAMERA 2 1 -0.9999701516465348 -0.0011120266840749639 -0.0075347911527510894 0.001298512589  
2 1 CAMERA 3
```

# Related Work

## Related Work

---

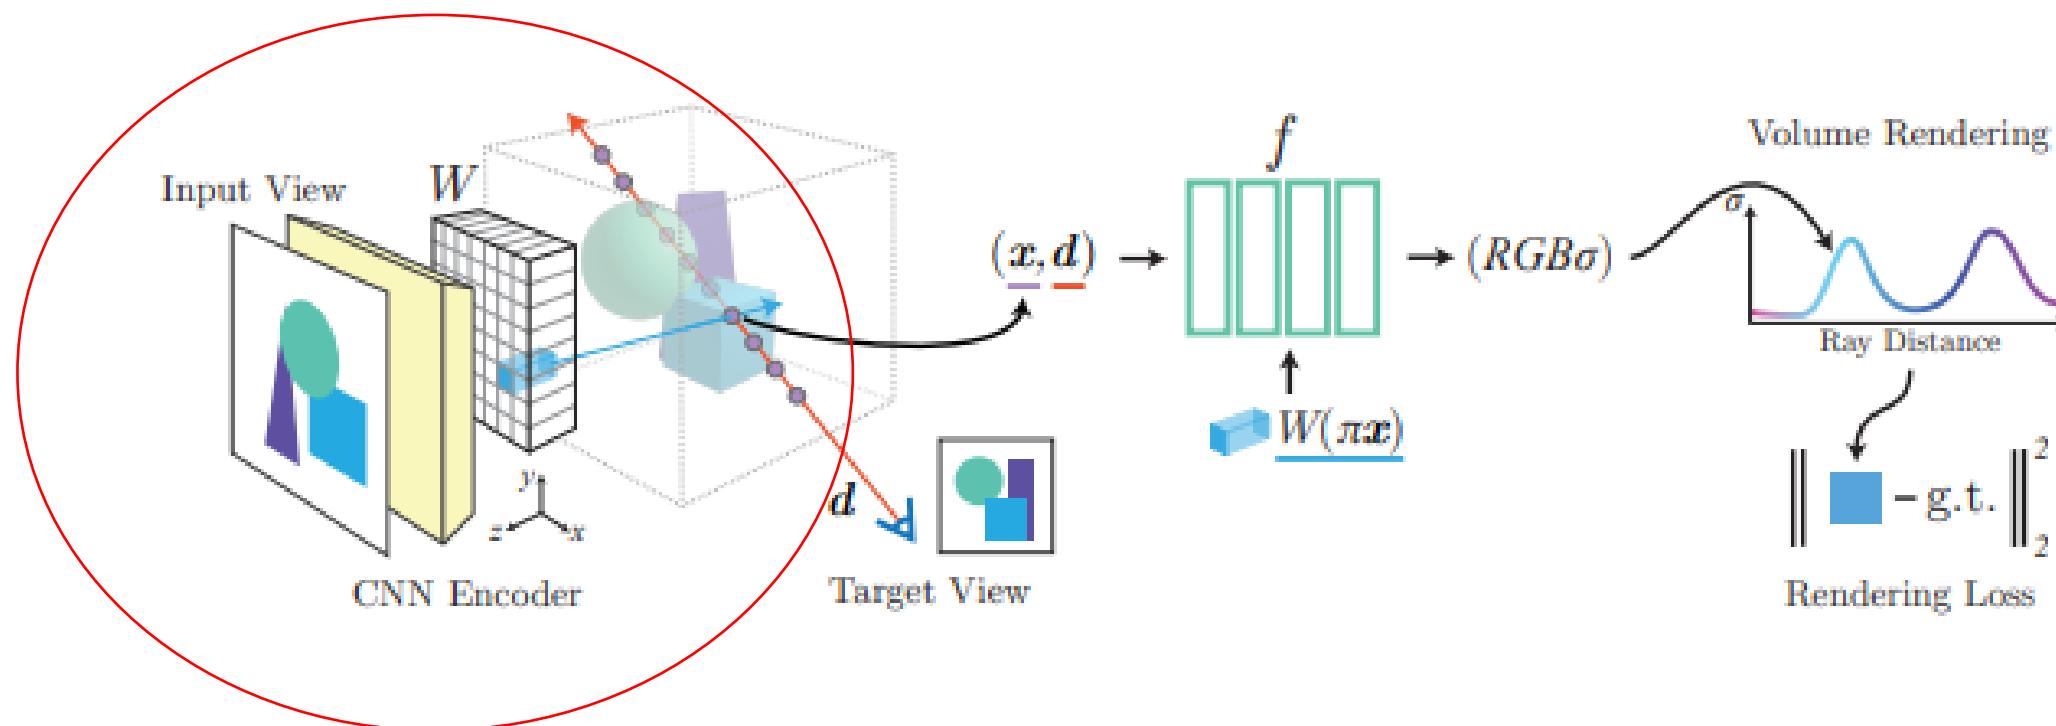
### PixelNeRF



## Related Work

---

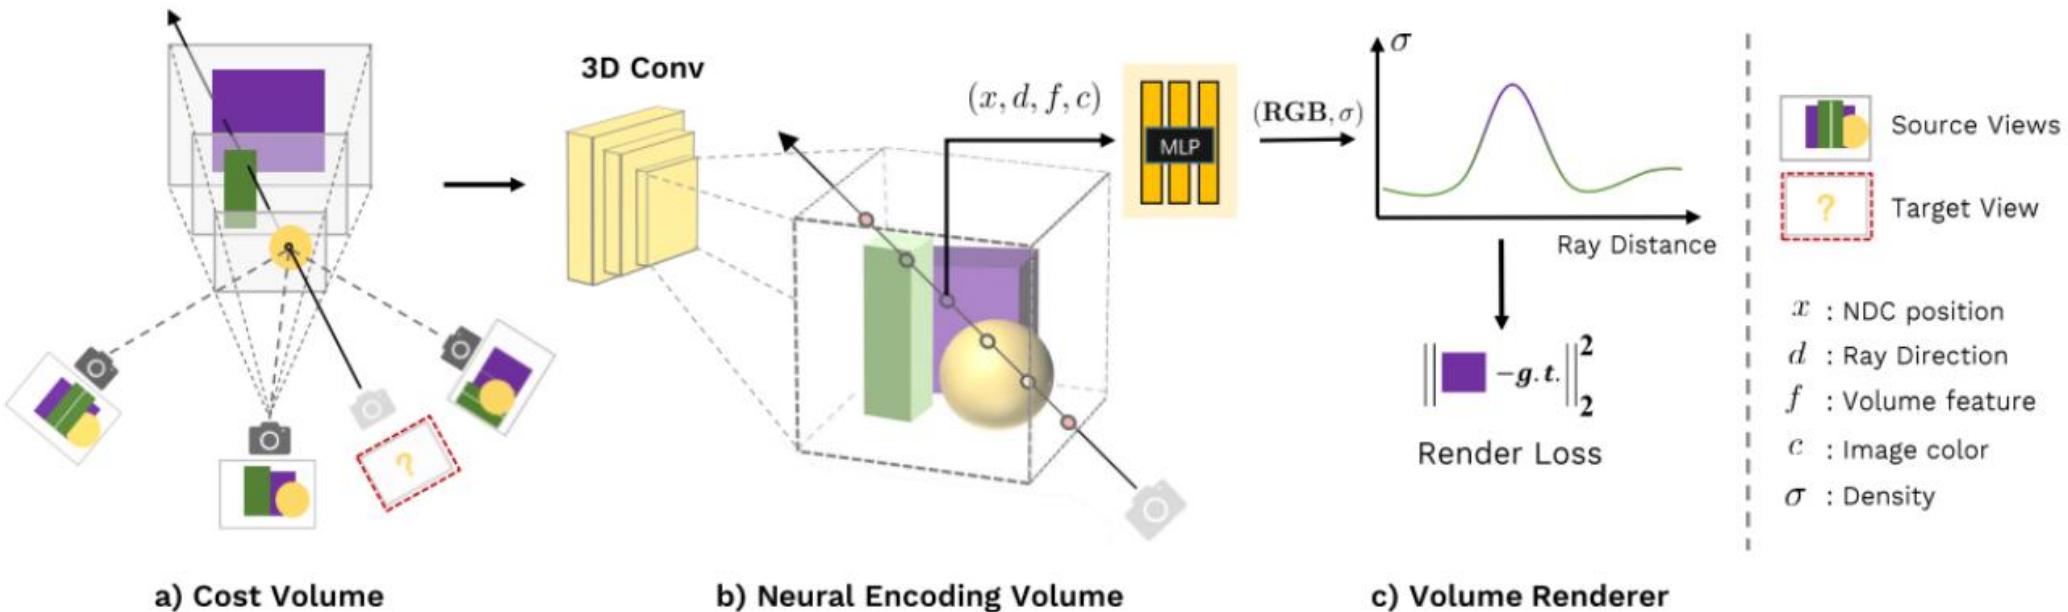
### PixelNeRF



# Related Work

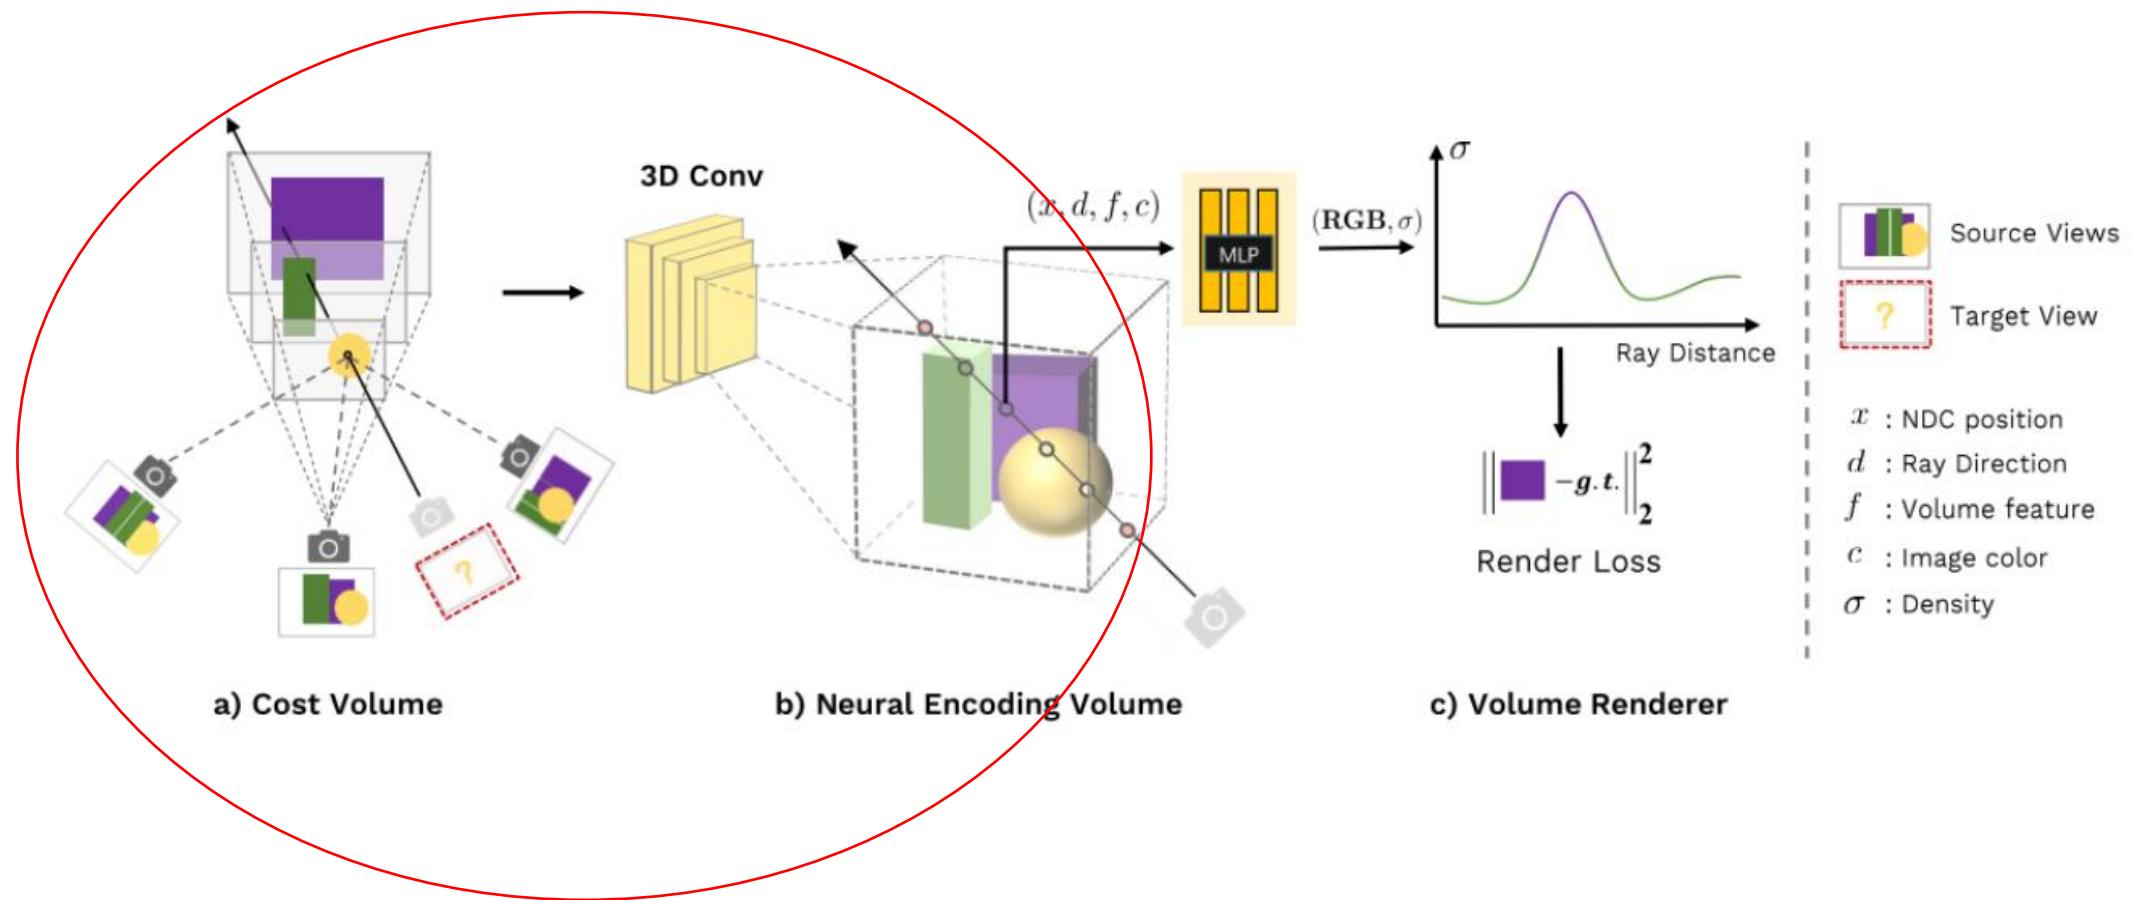
---

## MVS-NeRF



## Related Work

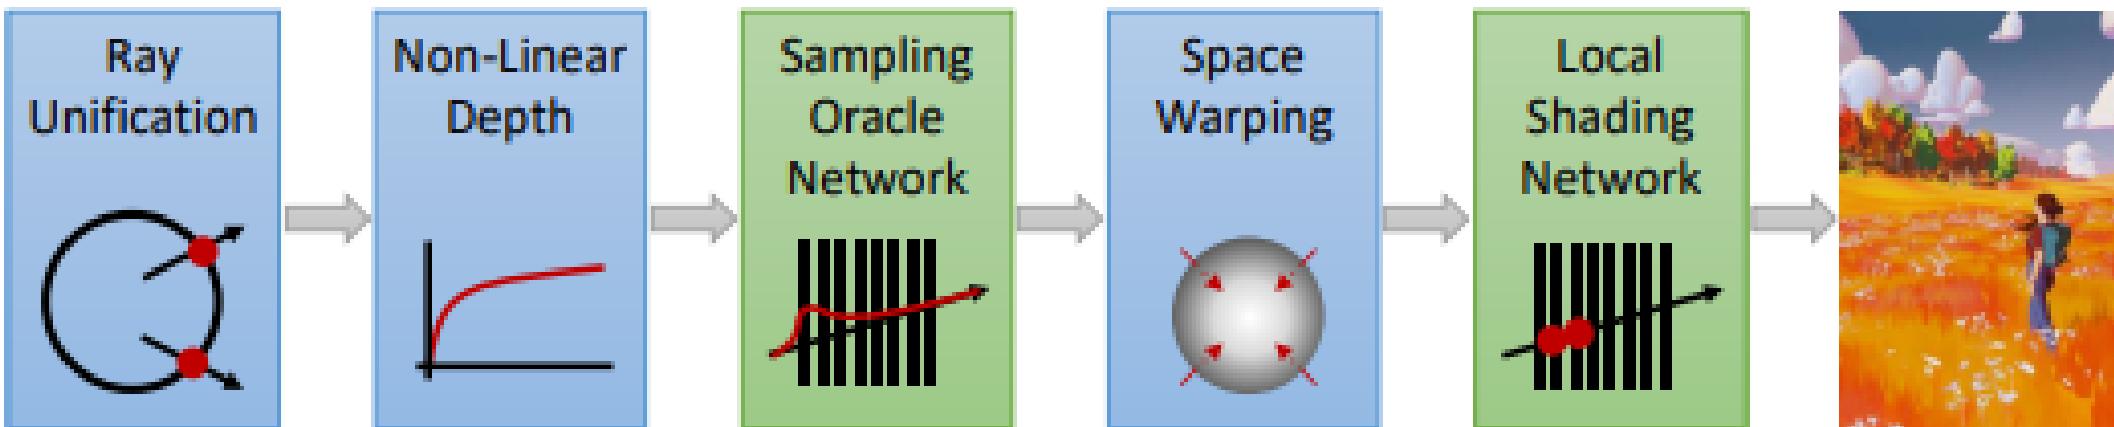
### MVS-NeRF



## Related Work

---

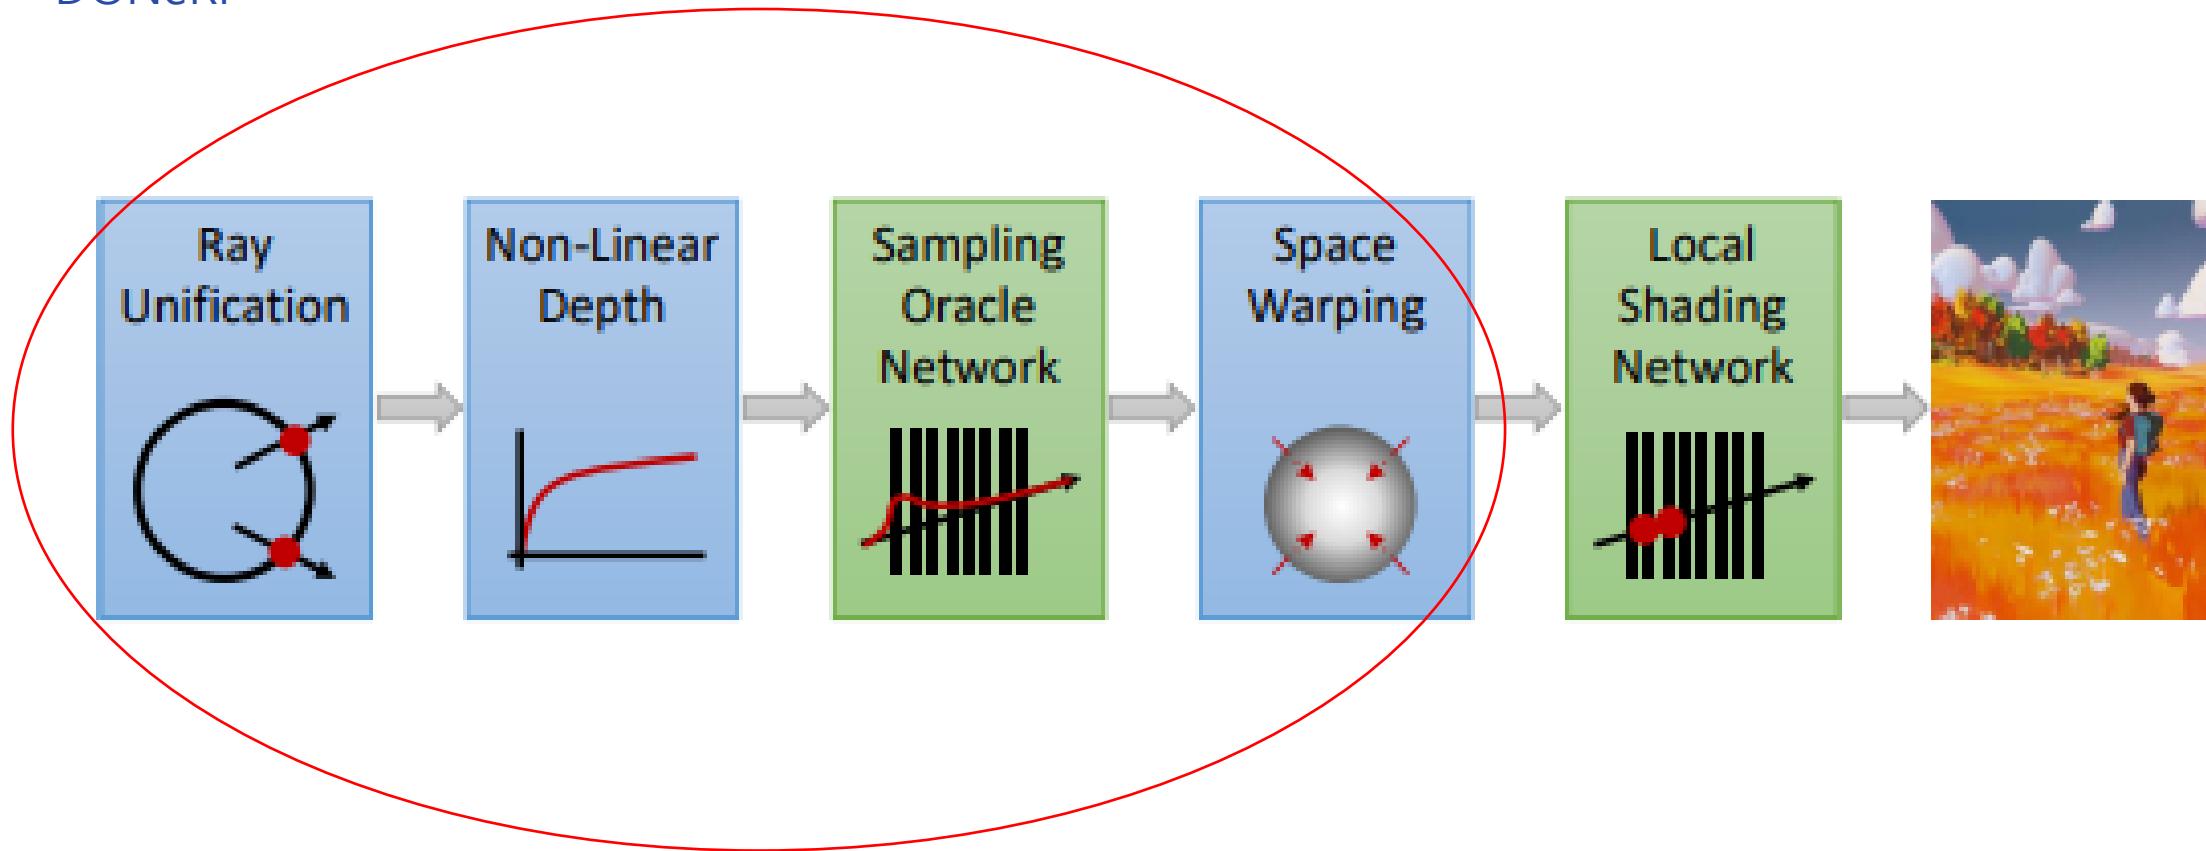
DONeRF



## Related Work

---

DONeRF



## Related Work

---

These are not Free.  
But, DS-NeRF is **Free**.

# Depth-Supervised Ray Termination

## Depth-Supervised Ray Termination

---

- DS-NeRF does **not** propose a **new network architecture**.
- It introduces an **additional** form of **supervision** to make the training of existing NeRF models more effective.
- Just add a new **L\_depth** to the **original Loss function**.

=> It can be **applied** to other **NeRFs**

$$\mathcal{L} = \mathcal{L}_{\text{Color}} + \lambda_D \mathcal{L}_{\text{Depth}}$$

## Depth-Supervised Ray Termination

---

$$\mathcal{L} = \underline{\mathcal{L}_{\text{Color}}} + \lambda_D \mathcal{L}_{\text{Depth}}$$

Same as NeRF

$$\hat{\mathbf{C}} = \int_0^{\infty} T(t) \sigma(t) \mathbf{c}(t) dt,$$

$$\hat{\mathbf{C}} = \int_0^{\infty} h(t) \mathbf{c}(t) dt = \mathbb{E}_{h(t)}[\mathbf{c}(t)].$$

$$\mathcal{L}_{\text{Color}} = \mathbb{E}_{\mathbf{r} \in \mathcal{R}(\mathbf{P})} \left\| \hat{\mathbf{C}}(\mathbf{r}) - \mathbf{C}(\mathbf{r}) \right\|_2^2.$$

## Depth-Supervised Ray Termination

---

NeRF Loss

$$\mathcal{L} = \sum_{\mathbf{r} \in \mathcal{R}} \left[ \left\| \hat{C}_c(\mathbf{r}) - C(\mathbf{r}) \right\|_2^2 + \left\| \hat{C}_f(\mathbf{r}) - C(\mathbf{r}) \right\|_2^2 \right]$$

Some Difference

DS-NeRF Color\_Loss

$$\mathcal{L}_{\text{Color}} = \mathbb{E}_{\mathbf{r} \in \mathcal{R}(\mathbf{P})} \left\| \hat{\mathbf{C}}(\mathbf{r}) - \mathbf{C}(\mathbf{r}) \right\|_2^2.$$

## Depth-Supervised Ray Termination

---

NeRF Loss

$$\mathcal{L} = \sum_{\mathbf{r} \in \mathcal{R}} \left[ \left\| \hat{C}_c(\mathbf{r}) - C(\mathbf{r}) \right\|_2^2 + \left\| \hat{C}_f(\mathbf{r}) - C(\mathbf{r}) \right\|_2^2 \right]$$

SUM

=

DS-NeRF Color\_Loss

$$\mathcal{L}_{\text{Color}} = \mathbb{E}_{\mathbf{r} \in \mathcal{R}(\mathbf{P})} \left\| \hat{\mathbf{C}}(\mathbf{r}) - \mathbf{C}(\mathbf{r}) \right\|_2^2.$$

Average

## Depth-Supervised Ray Termination

---

$$\mathcal{L} = \mathcal{L}_{\text{Color}} + \lambda_D \mathcal{L}_{\text{Depth}}$$

hyper-parameter

our objective is to minimize the KL divergence between the rendered ray distribution  $h_{ij}(t)$  of  $\mathbf{x}_i$ 's image coordinates and the noisy depth distribution:

## Depth-Supervised Ray Termination

---

$$\mathcal{L} = \mathcal{L}_{\text{Color}} + \lambda_D \underline{\mathcal{L}_{\text{Depth}}}$$

our objective is to minimize the KL divergence between the rendered ray distribution  $h_{ij}(t)$  of  $\mathbf{x}_i$ 's image coordinates and the noisy depth distribution:  $h(t) = T(t)\sigma(t)$

$$\mathbb{N}(\mathbf{D}_{ij}, \hat{\sigma}_i)$$

Sparse 3D Point Cloud(의 Z값 = Depth)

points3D.txt	
This file contains the information of all reconstructed 3D points in the dataset using one line per point, e.g.:	
<b>재투영오차</b>	
# 3D point list with one line of data per point:	# POINT3D_ID, X, Y, Z, R, G, B, [ERROR, BACK]
# Number of points: 3, mean track length: 3.3334	63398 1.67248 0.292931 0.699726 115 121 122 1.33927 16 6542 15 7345 6 6714 14 7227 63376 2.01848 0.198877 -0.8268841 182 209 250 1.73449 16 6519 15 7322 14 7212 8 3991 63371 1.71182 0.28566 0.53475 245 251 249 0.612829 118 4148 117 4473

## Depth-Supervised Ray Termination

---

$$\mathbb{E}_{\mathbb{D}_{ij}} \text{KL}[\delta(t - \mathbb{D}_{ij}) || h_{ij}(t)] = \text{KL}[\mathcal{N}(\mathbf{D}_{ij}, \hat{\sigma}_i) || h_{ij}(t)] + \text{const.}$$

i : 3D KeyPoint Number  
j: Camera(image) Number

Using definition of KL Divergence

$$\text{KL}[p||q] = \int p(x) \log \frac{p(x)}{q(x)} dx$$

$$\begin{aligned}\mathcal{L}_{Depth} &= \mathbb{E}_{x_i \in X_j} \left[ - \int \log h(t) \exp \left( -\frac{(t - \mathbf{D}_{ij})^2}{2\hat{\sigma}_i^2} \right) dt \right] \\ &\approx \mathbb{E}_{x_i \in X_j} \left[ - \sum_k \log h_k \exp \left( -\frac{(t_k - \mathbf{D}_{ij})^2}{2\hat{\sigma}_i^2} \right) \Delta t_k \right].\end{aligned}$$

# Experiment

# Experiment

---

DataSet

**DTU MVS Dataset**



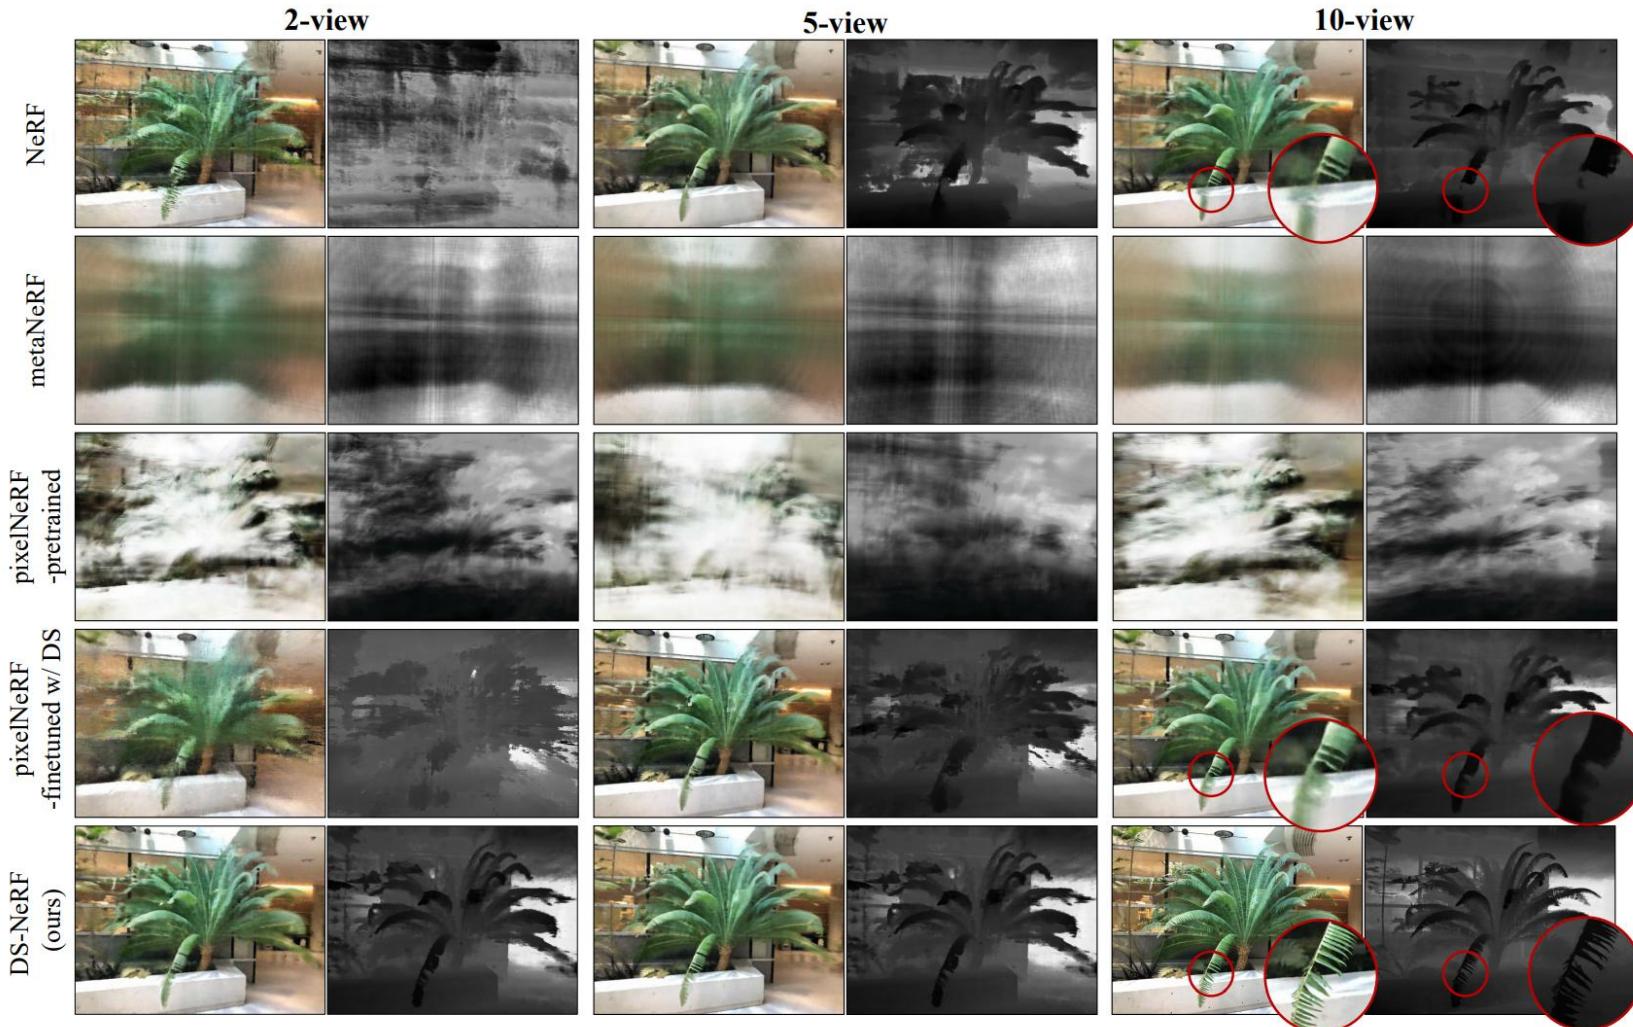
**NeRF Real-world Data**

**Redwood-3dscan**



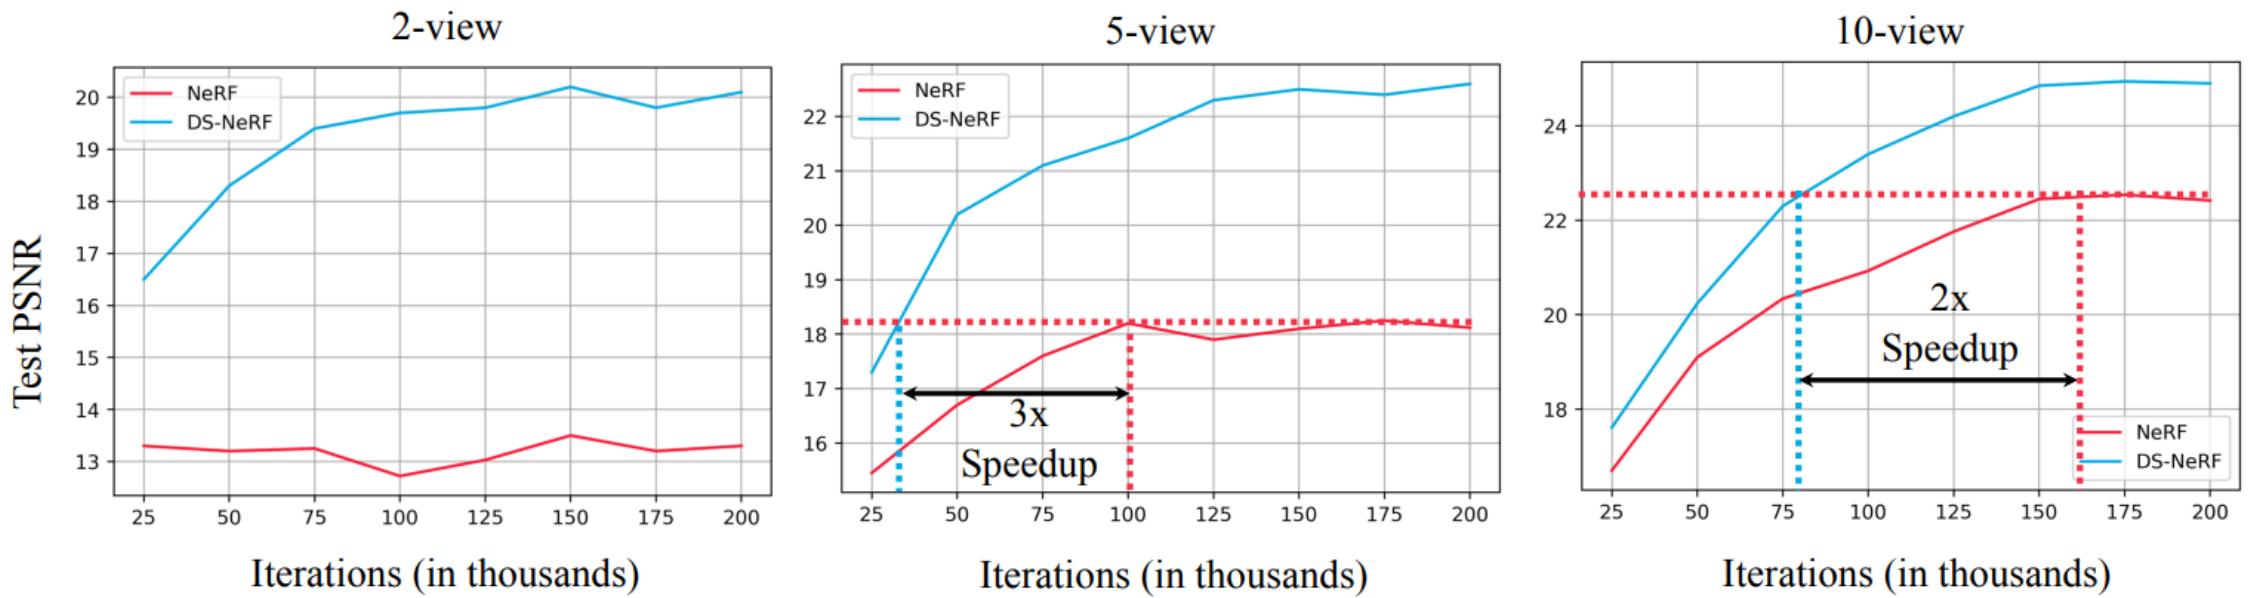
# Experiment

---



# Experiment

---



2~3 times faster

## Experiment

---

DTU [8]	PSNR↑			SSIM↑			LPIPS↓		
	3-view	6-view	9-view	3-view	6-view	9-view	3-view	6-view	9-view
NeRF	9.9	18.6	22.1	0.37	0.72	<b>0.82</b>	0.62	0.35	0.26
metaNeRF-DTU	18.2	18.8	20.2	0.60	0.61	0.67	0.40	0.41	0.35
pixelNeRF-DTU	<b>19.3</b>	20.4	21.1	<b>0.70</b>	0.73	0.76	<b>0.39</b>	0.36	0.34
DS-NeRF									
MSE	16.5	20.5	22.2	0.54	0.73	0.77	0.48	0.31	0.26
KL divergence	16.9	<b>20.6</b>	<b>22.3</b>	0.57	<b>0.75</b>	0.81	0.45	<b>0.29</b>	<b>0.24</b>

## Contribution

---

- Addressing NeRF's Limitations: Fewer Views & Faster Training
- Proposing a Depth Supervision Loss for Free
- High Compatibility