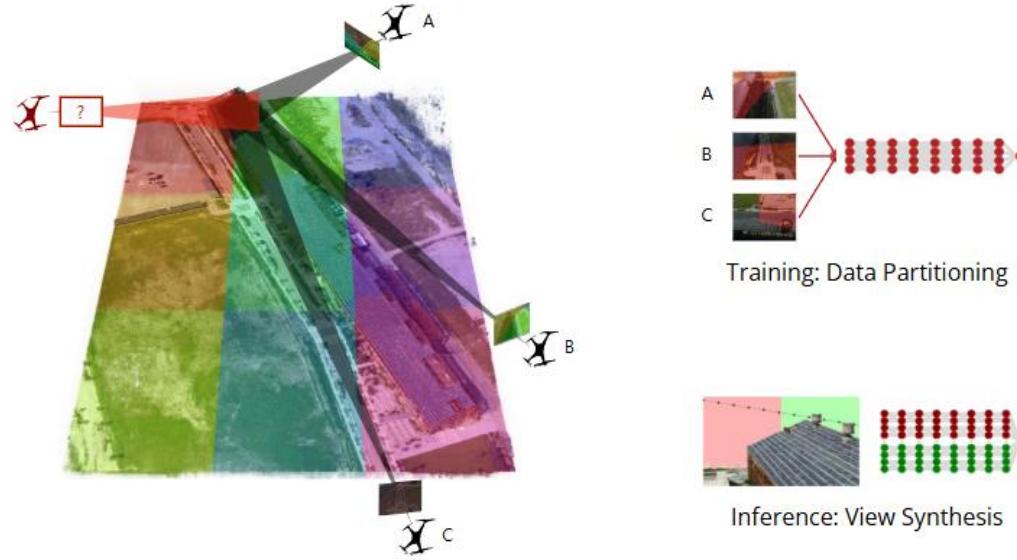


Mega-NeRF: Scalable Construction of Large-Scale NeRFs for Virtual Fly-Throughs (CVPR 2022)



한국과학기술연구원(KIST)

CVIPL 학생연구원 김연욱

Contents

- Author
- OverView
- Introduction
- Related Work
- Method
- Experiment
- Contribution

Author

Mega-NeRF: Scalable Construction of Large-Scale NeRFs for Virtual Fly-Throughs

Haithem Turki¹

Deva Ramanan^{1,2}

¹Carnegie Mellon University

Mahadev Satyanarayanan¹

³Argo AI



Deva Ramanan

Professor, Robotics Institute, [Carnegie Mellon University](#).
cs.cmu.edu의 이메일 확인됨 - [홈페이지](#)

Computer Vision Machine Learning

팔로우

제목	인용	연도
Microsoft coco: Common objects in context TY Lin, M Maire, S Belongie, J Hays, P Perona, D Ramanan, P Dollár, ... European conference on computer vision, 740-755	63300	2014
Object detection with discriminatively trained part-based models PF Felzenszwalb, RB Girshick, D McAllester, D Ramanan IEEE transactions on pattern analysis and machine intelligence 32 (9), 1627-1645	13643	2009
A discriminatively trained, multiscale, deformable part model P Felzenszwalb, D McAllester, D Ramanan 2008 IEEE conference on computer vision and pattern recognition, 1-8	4164	2008
Depth-supervised nerf: Fewer views and faster training for free K Deng, A Liu, JY Zhu, D Ramanan Proceedings of the IEEE/CVF conference on computer vision and pattern ...	1141	2022

Depth-Supervised Nerf
나온 이후 나온 논문

Mega-NeRF: Scalable Construction of Large-Scale NeRFs for Virtual Fly-Throughs

Haithem Turki¹

¹Carnegie Mellon University

Deva Ramanan^{1,2}

Mahadev Satyanarayanan¹

²Argo AI



Haithem Turki

Research Scientist, [NVIDIA](#)
nvidia.com의 이메일 확인됨 - [홈페이지](#)

팔로우



Mahadev Satyanarayanan

Jaime Carbonell University Professor of Computer Science, [Carnegie Mellon University](#)
cs.cmu.edu의 이메일 확인됨 - [홈페이지](#)

팔로우

edge computing mobile computing Internet of Things pervasive computing distributed systems

제목

인용

연도

[Mega-nerf: Scalable construction of large-scale nerfs for virtual fly-throughs](#)

488

2022

H Turki, D Ramanan, M Satyanarayanan
Proceedings of the IEEE/CVF conference on computer vision and pattern ...

[Suds: Scalable urban dynamic scenes](#)

152

2023

H Turki, JY Zhang, F Ferroni, D Ramanan
Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern ...

[Learning specific-class segmentation from diverse data](#)

112

2011

MP Kumar, H Turki, D Preston, D Koller
2011 International conference on computer vision, 1800-1807

제목

인용

연도

[The case for vm-based cloudlets in mobile computing](#)

4994

2009

M Satyanarayanan, P Bahl, R Caceres, N Davies
IEEE pervasive Computing 8 (4), 14-23

[Pervasive computing: Vision and challenges](#)

4015

2002

M Satyanarayanan
IEEE Personal communications 8 (4), 10-17

[The emergence of edge computing](#)

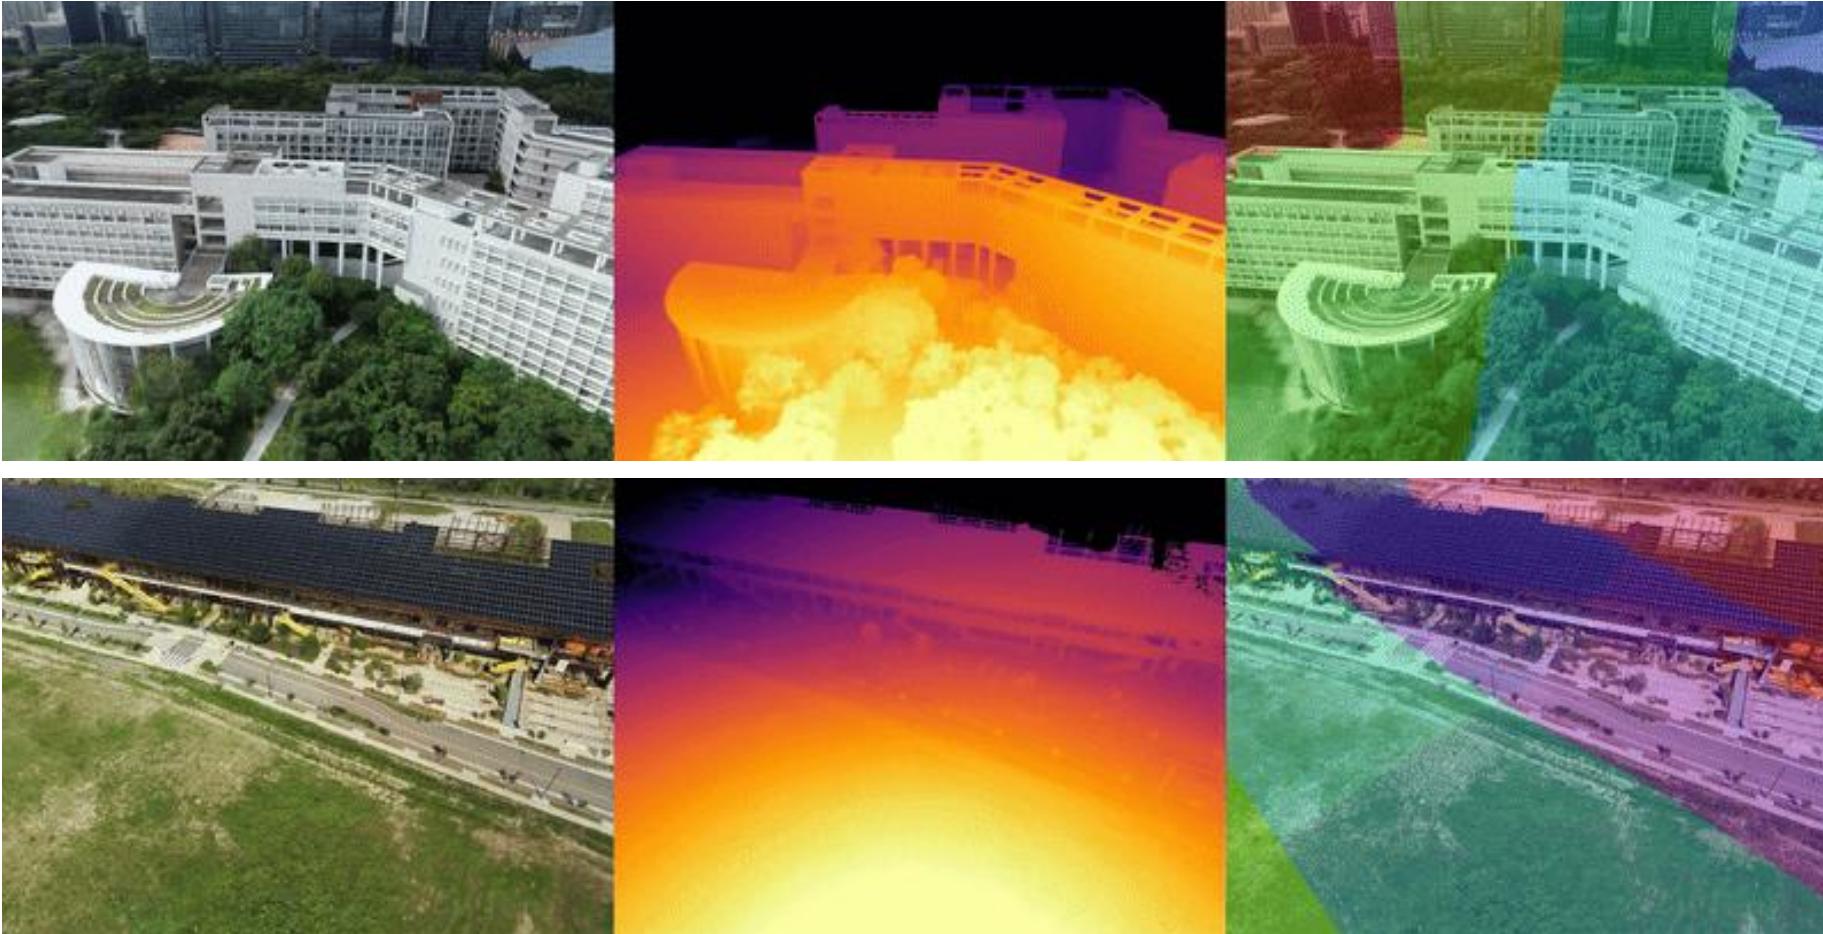
3245

2017

M Satyanarayanan
Computer 50 (1), 30-39

OverView

OverView

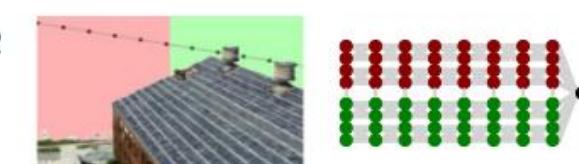
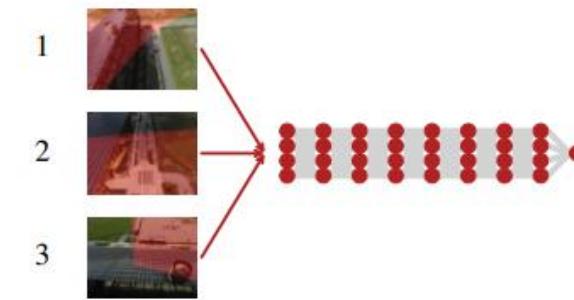
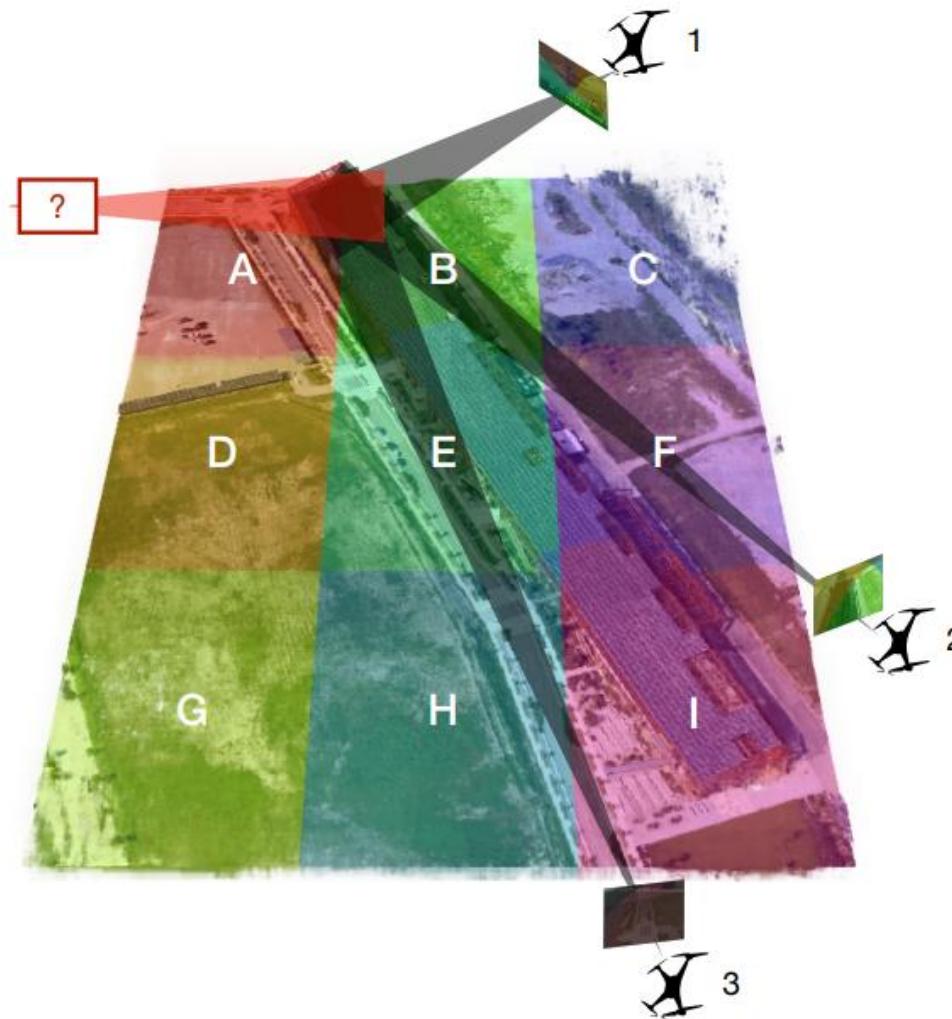


Reconstructed Image

Depth Map

Submodule Visualization

OverView



Introduction

Introduction

Search and rescue (SAR) :

search for and provision of aid to people who are in distress or imminent danger.

기존 Drone based SAR 문제점:

- 배터리 부족
- 2D "Birds-eye-view" map만 생성
--> 디테일 정보 부족



2D "Birds-eye-view" map

Introduction

Search and rescue (SAR) :

search for and provision of aid to people
who are in distress or imminent danger.

기존 Drone based SAR 문제점:

- 배터리 부족
- 2D "Birds-eye-view" map만 생성
--> 디테일 정보 부족

NeRF로 해결!



2D "Birds-eye-view" map

Introduction

- NeRF의 굉장히 많은 후속 연구
- 그러나, 대부분은 실내에서 찍은 **single-object scenes**에 대한 연구
- NeRF의 가장 큰 데이터셋 Tanks and Temples Dataset의 평균 촬영 공간의 넓이는 Scene 당 $463m^2$



Introduction

- 저자가 원하는 것은 Urban-Scale environments NeRF
- Urban-Scale environments의 Dataset은 order-of-magnitude more pixels을 포함한다.

	Resolution	# Images	# Pixels/Rays	Scene Captured / Image
Synthetic NeRF - Chair	400 x 400	400	256,000,000	0.271
Synthetic NeRF - Drums	400 x 400	400	256,000,000	0.302
Synthetic NeRF - Ficus	400 x 400	400	256,000,000	0.582
Synthetic NeRF - Hotdog	400 x 400	400	256,000,000	0.375
Synthetic NeRF - Lego	400 x 400	400	256,000,000	0.205
Synthetic NeRF - Materials	400 x 400	400	256,000,000	0.379
Synthetic NeRF - Mic	400 x 400	400	256,000,000	0.518
Synthetic NeRF - Ship	400 x 400	400	256,000,000	0.483
T&T - Barn	1920 x 1080	384	796,262,400	0.135
T&T - Caterpillar	1920 x 1080	368	763,084,800	0.216
T&T - Family	1920 x 1080	152	315,187,200	0.284
T&T - Ignatius	1920 x 1080	263	545,356,800	0.476
T&T - Truck	1920 x 1080	250	518,400,000	0.225
Mill 19 - Building	4608 x 3456	1940	30,894,981,120	0.062
Mill 19 - Rubble	4608 x 3456	1678	26,722,566,144	0.050
Quad 6k	1708 x 1329	5147	11,574,265,679	0.010
UrbanScene3D - Residence	5472 x 3648	2582	51,541,512,192	0.059
UrbanScene3D - Sci-Art	4864 x 3648	3019	53,568,749,568	0.088
UrbanScene3D - Campus	5472 x 3648	5871	117,196,056,576	0.028



Mega-Nerf Dataset

Introduction

- Firstly, applications such as search-and-rescue are **time-sensitive**.
- Secondly, our **datasets** are **orders of magnitude larger** than previously evaluated datasets.
- Finally, existing real-time NeRF renderers are **ill-suited** for large-scale scenes

Introduction

- Firstly, applications such as search-and-rescue are **time-sensitive**.
- Secondly, our **datasets** are **orders of magnitude larger** than previously evaluated datasets.

해결방법 :

거대한 전체 장면을 여러 조각(Cell)으로 나누어 각 조각(Cell)을 submodules(NeRF)이 전담

+

submodules 병렬로 훈련

Introduction

해결방법 :

spatial locality라는 것을 사용

(설명은 뒤쪽에...)

- Finally, existing real-time NeRF renderers are **ill-suited** for large-scale scenes

Related Work

Related Work

DeRF

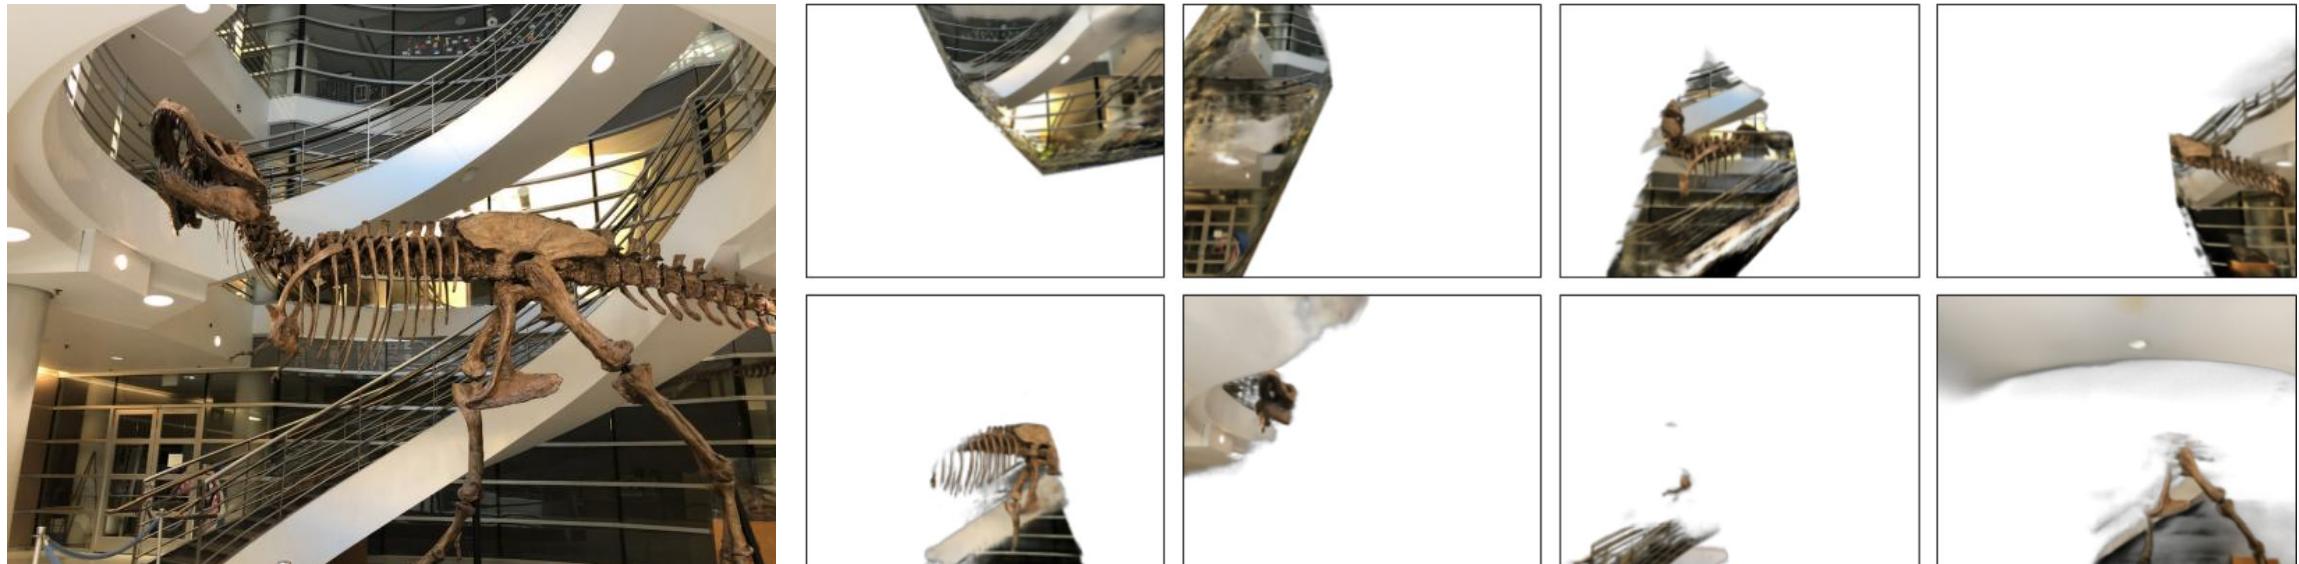


Figure 4. **Decomposed radiance fields** – We visualize each of the rendering heads individually. Note that as each head is rendered *only* the weights of *one* neural network head needs to be loaded, hence resulting in optimal cache coherency while accessing GPU memory.

Related Work

DeRF

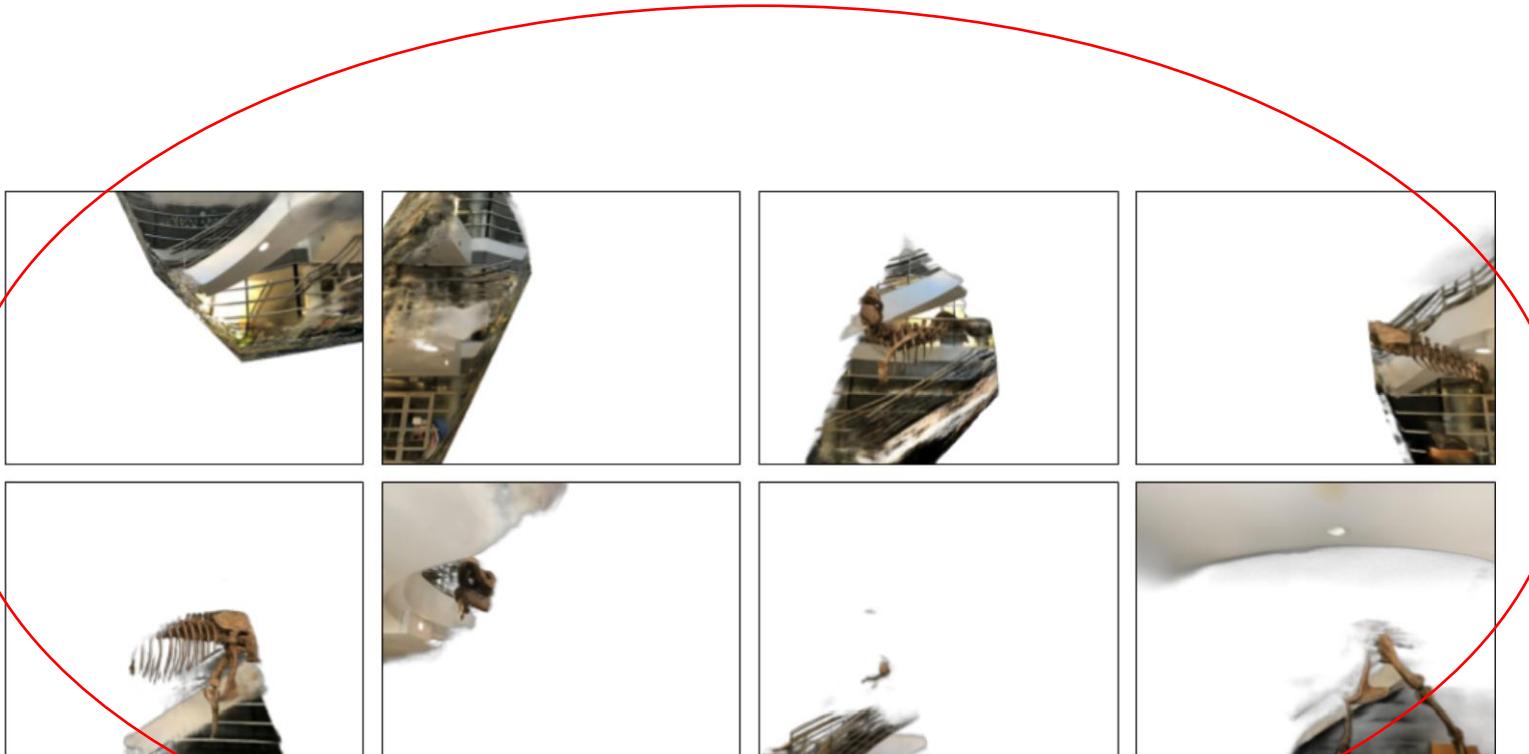


Figure 4. **Decomposed radiance fields** – We visualize each of the rendering heads individually. Note that as each head is rendered *only* the weights of *one* neural network head needs to be loaded, hence resulting in optimal cache coherency while accessing GPU memory.

Related Work

DeRF



“여러 개의 독립적인 mlp로 렌더링한다.” 아이디어 차용
spatial Voronoi partitioning을 통해
3D 장면을 여러 개의 독립적인 mlp로
나누어서 렌더링
→

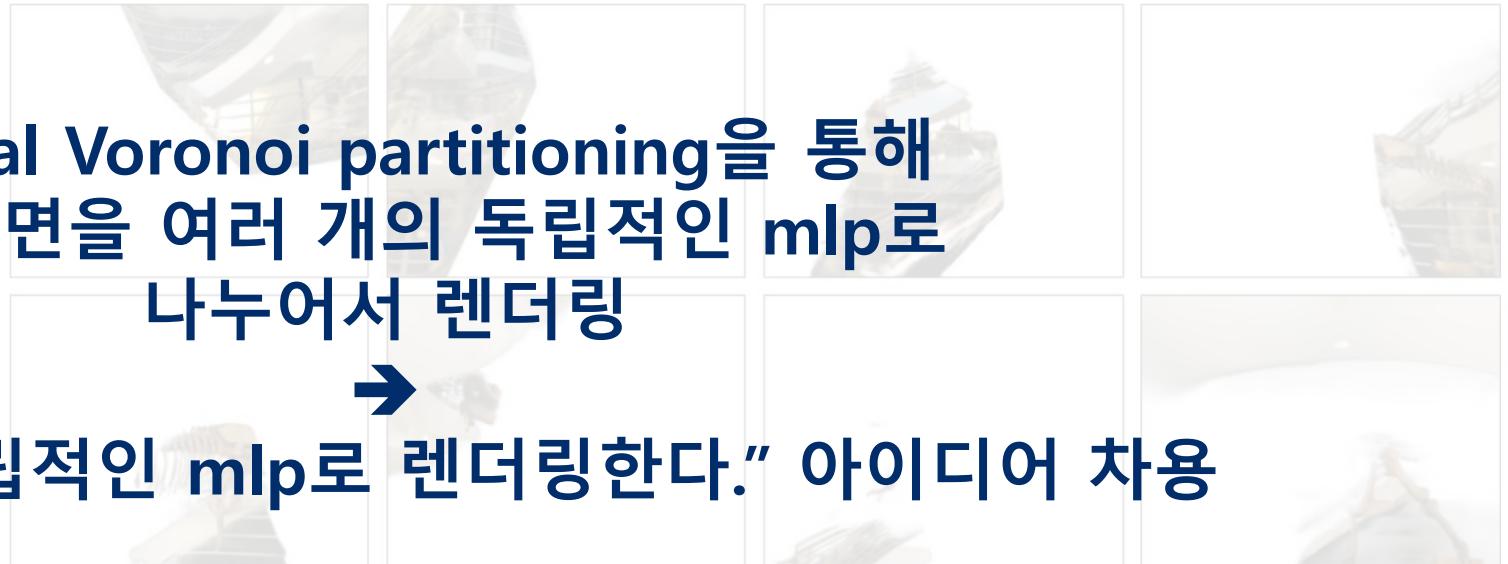
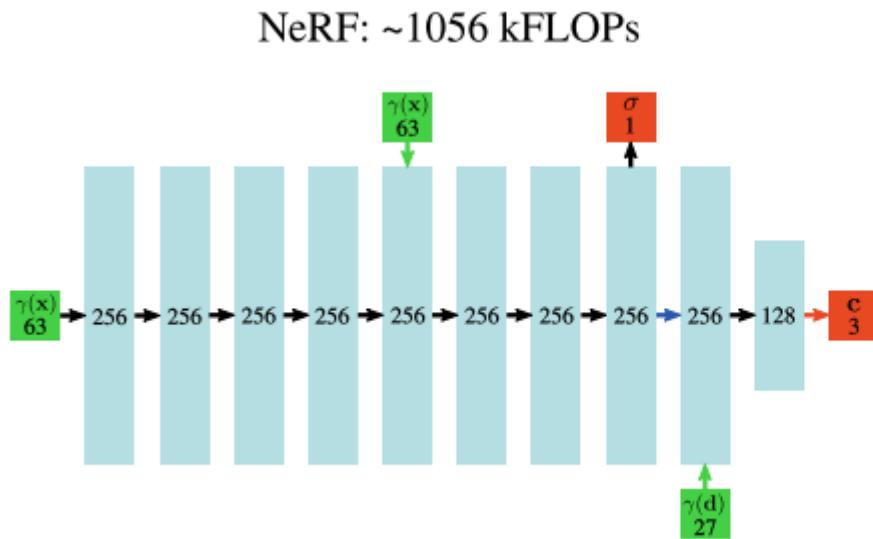


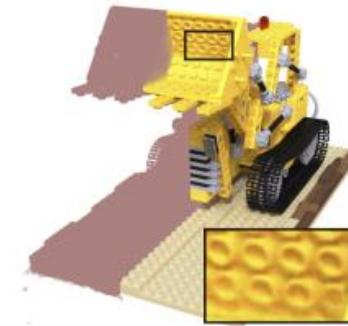
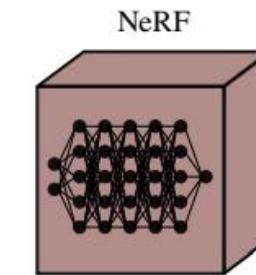
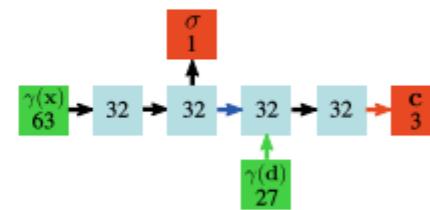
Figure 4. Decomposed radiance fields – We visualize each of the rendering heads individually. Note that as each head is rendered *only* the weights of *one* neural network head needs to be loaded, hence resulting in optimal cache coherency while accessing GPU memory.

Related Work

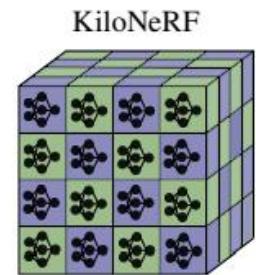
KiloNeRF



KiloNeRF: ~12 kFLOPs



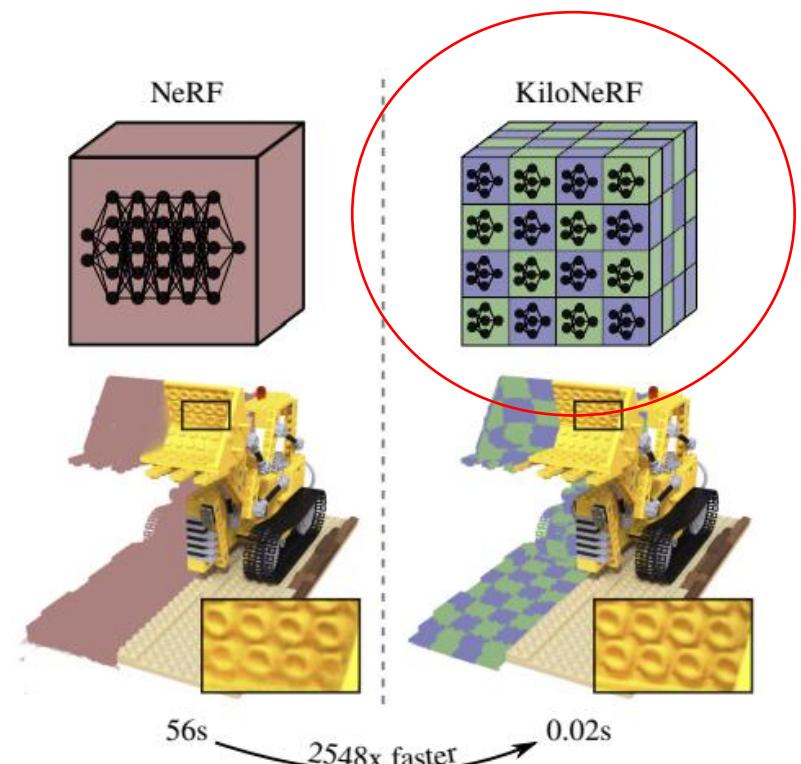
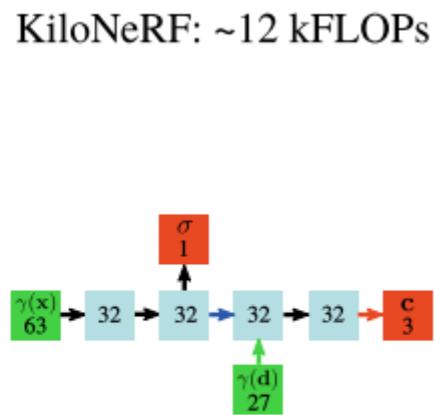
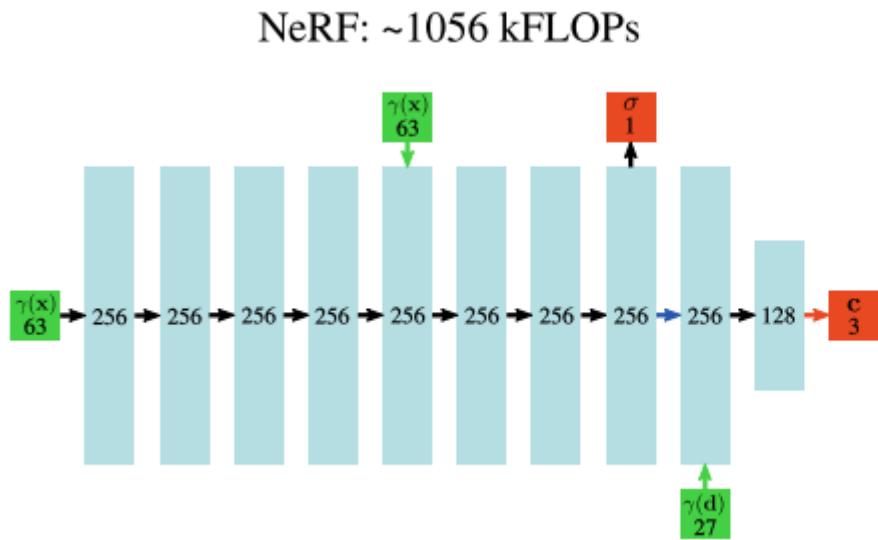
56s → 0.02s
2548x faster



56s → 0.02s
2548x faster

Related Work

KiloNeRF



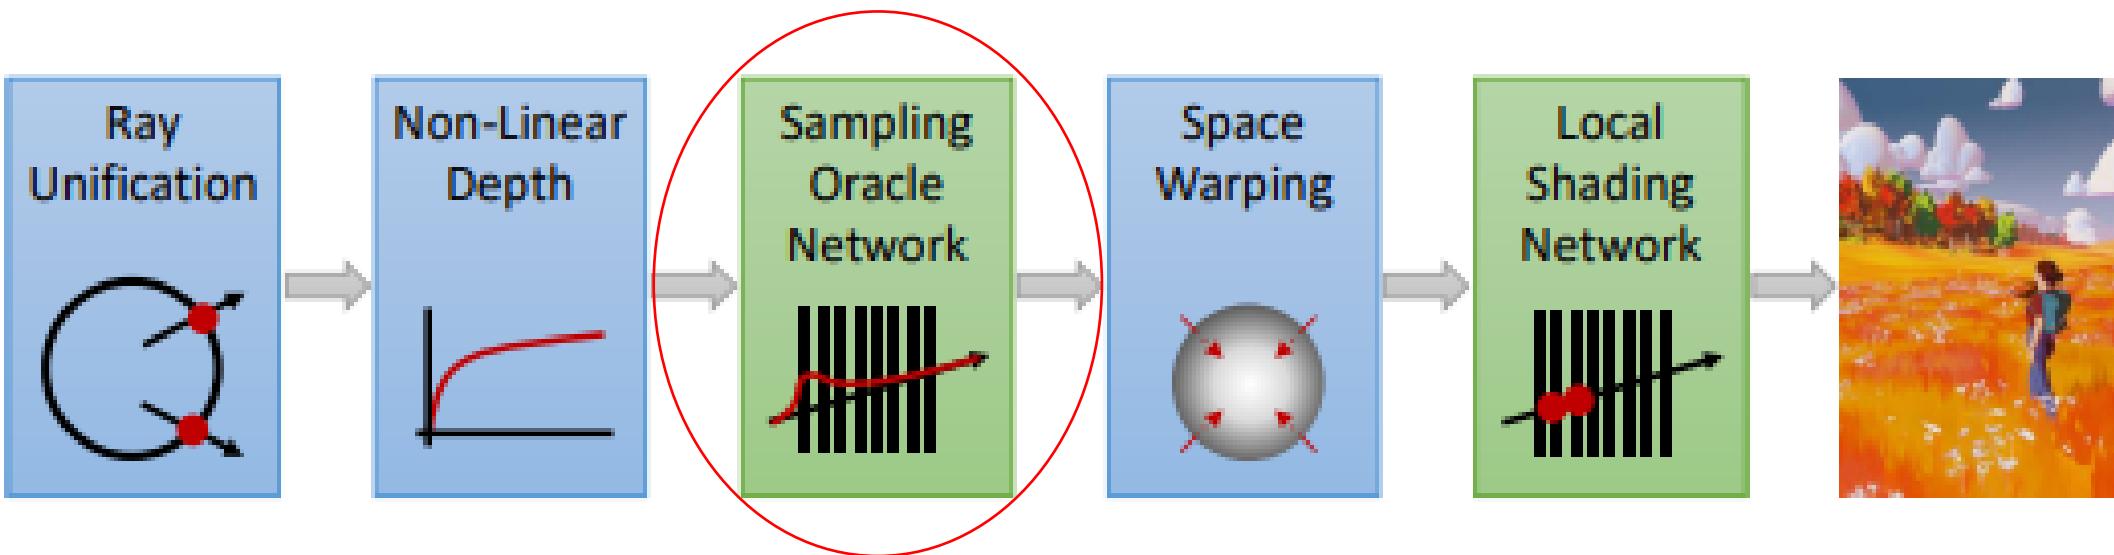
Related Work

KiloNeRF



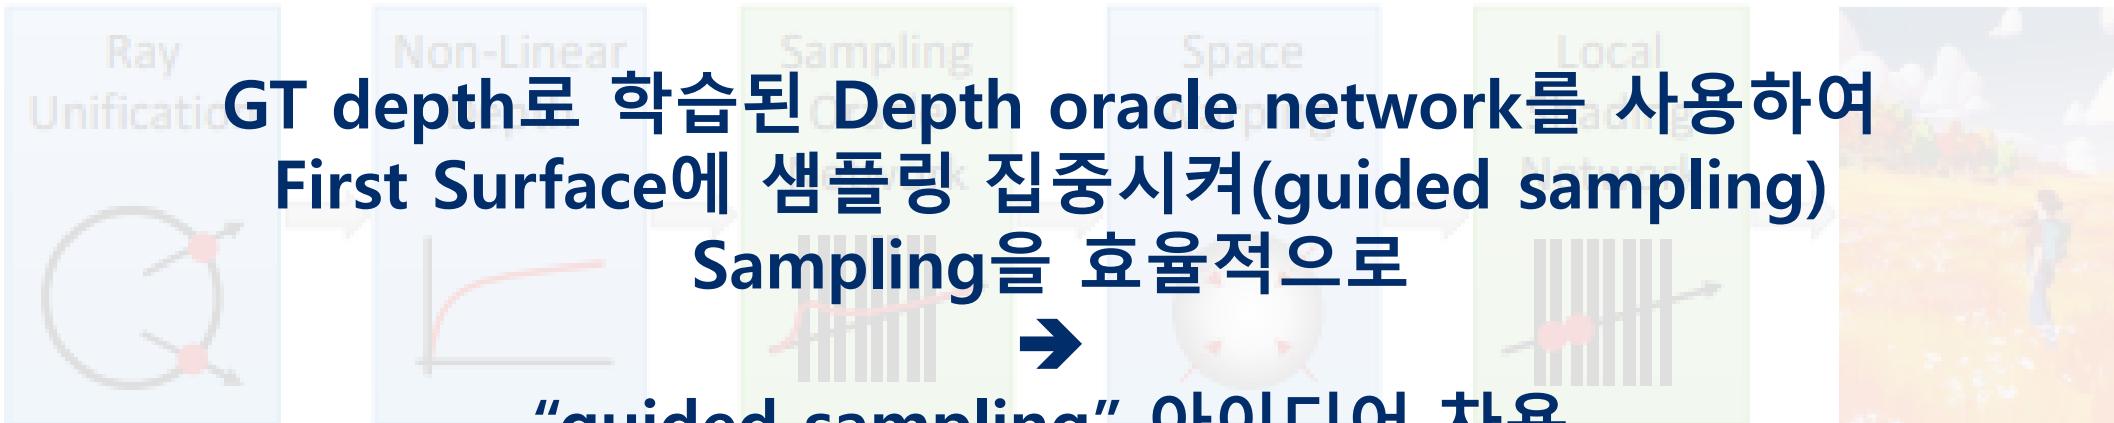
Related Work

DONeRF



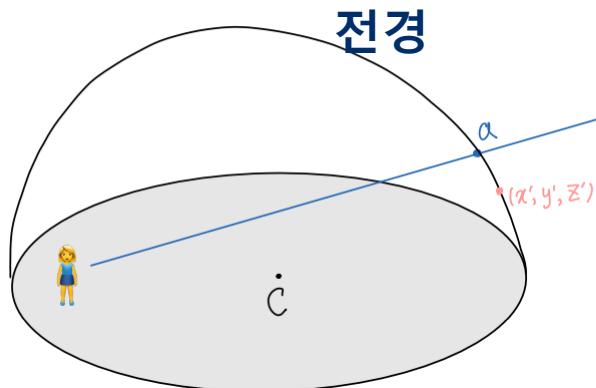
Related Work

DONeRF

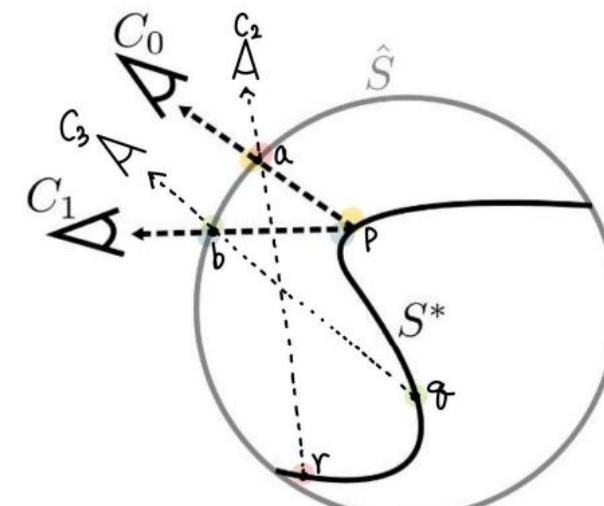


Related Work

NeRF++



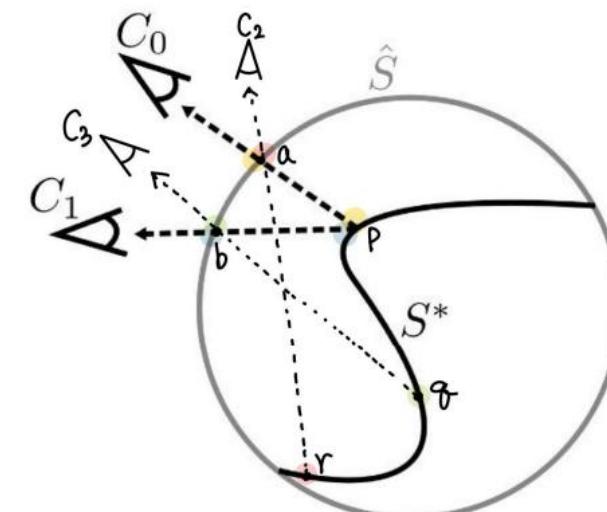
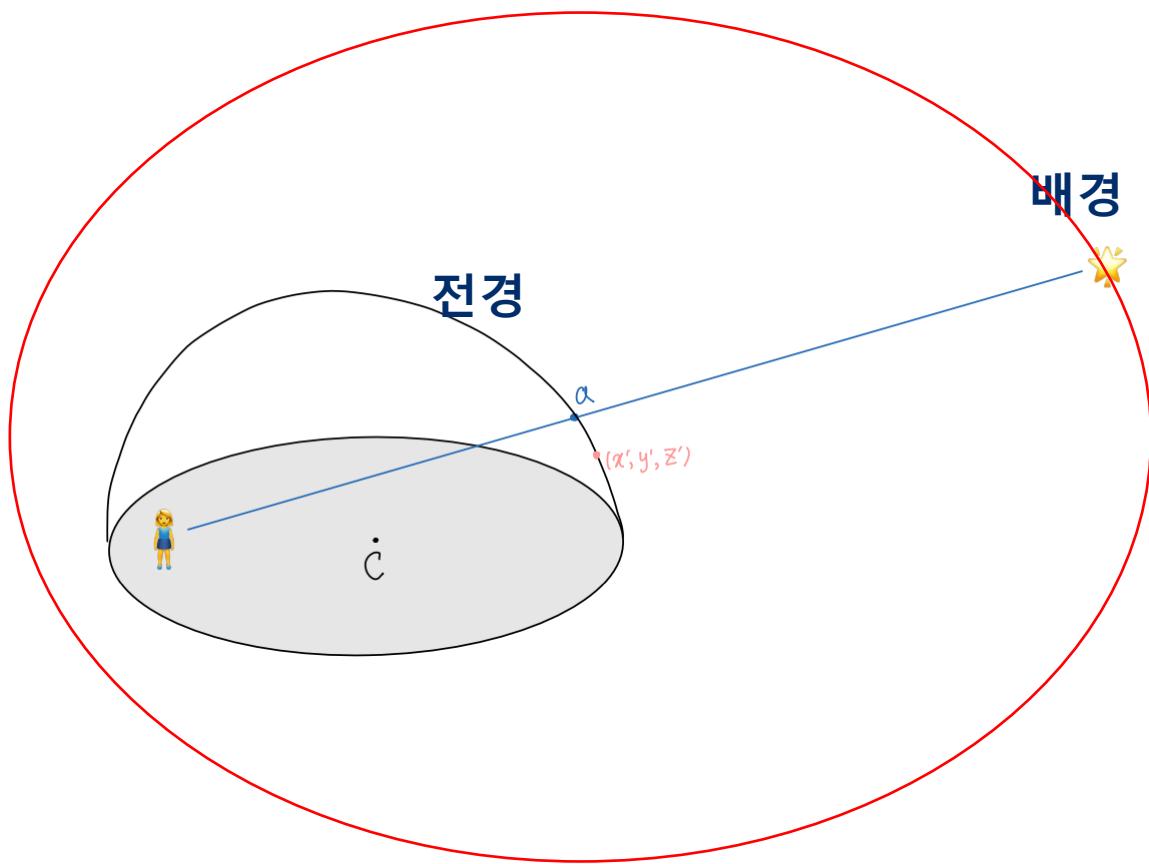
배경



Shape-Radiance Ambiguity

Related Work

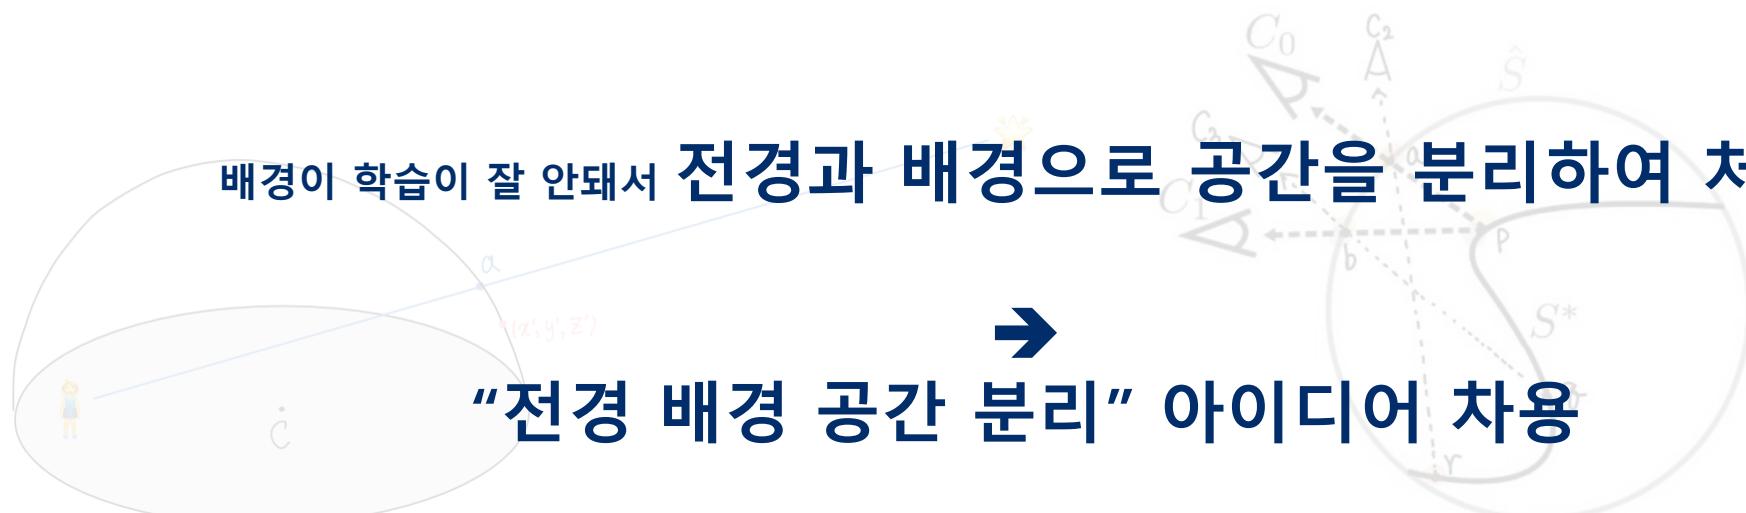
NeRF++



Shape-Radiance Ambiguity

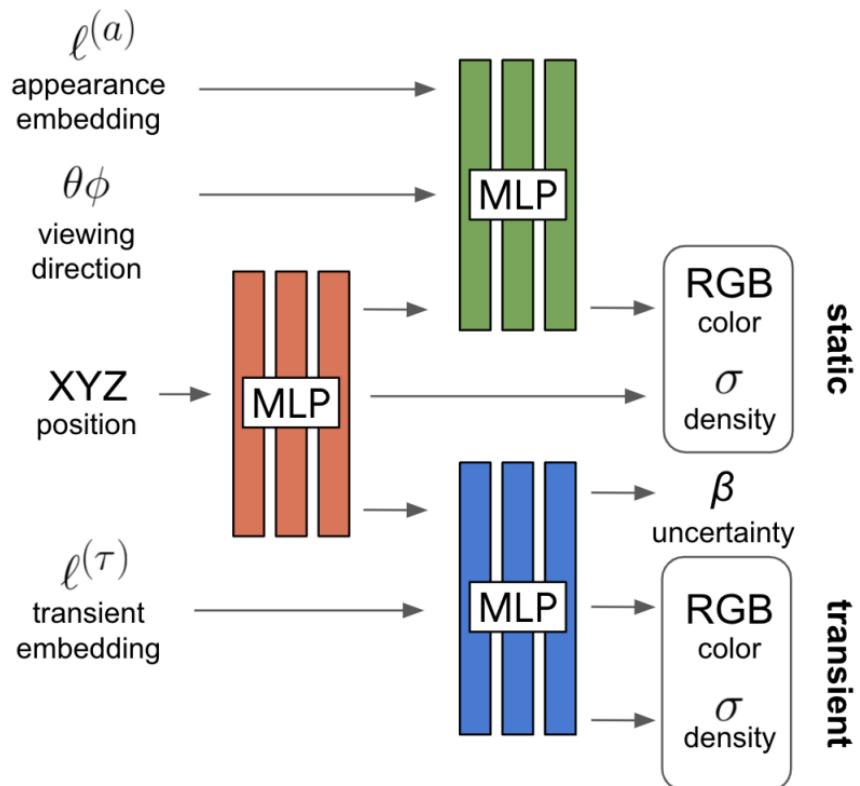
Related Work

NeRF++



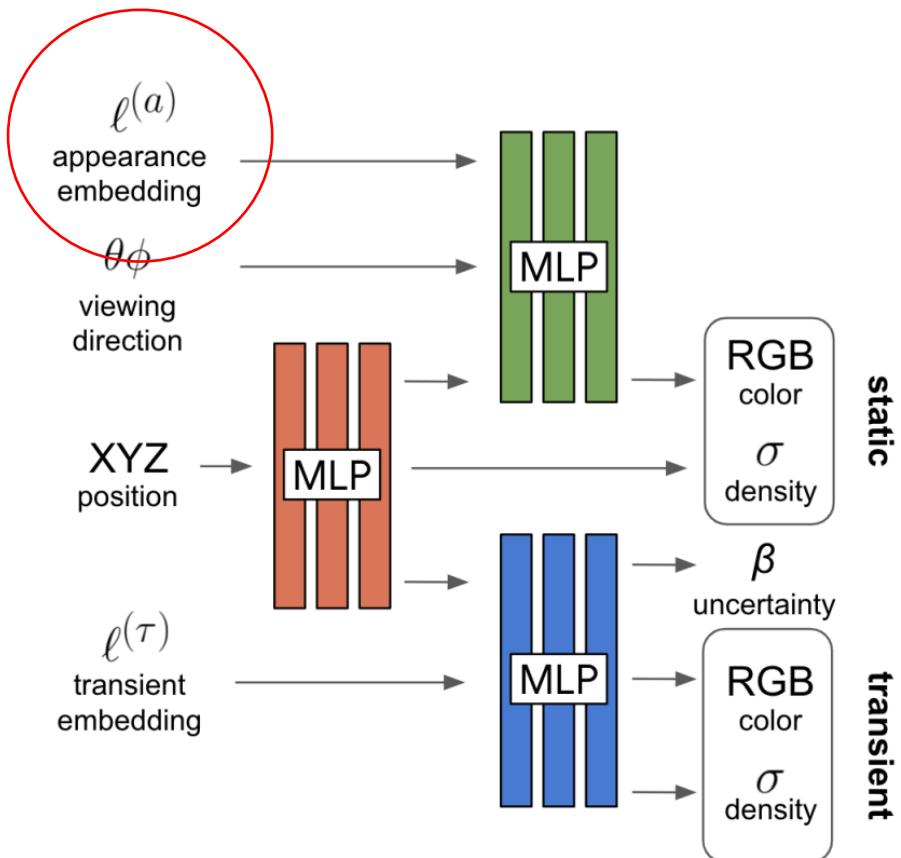
Related Work

NeRF in the Wild



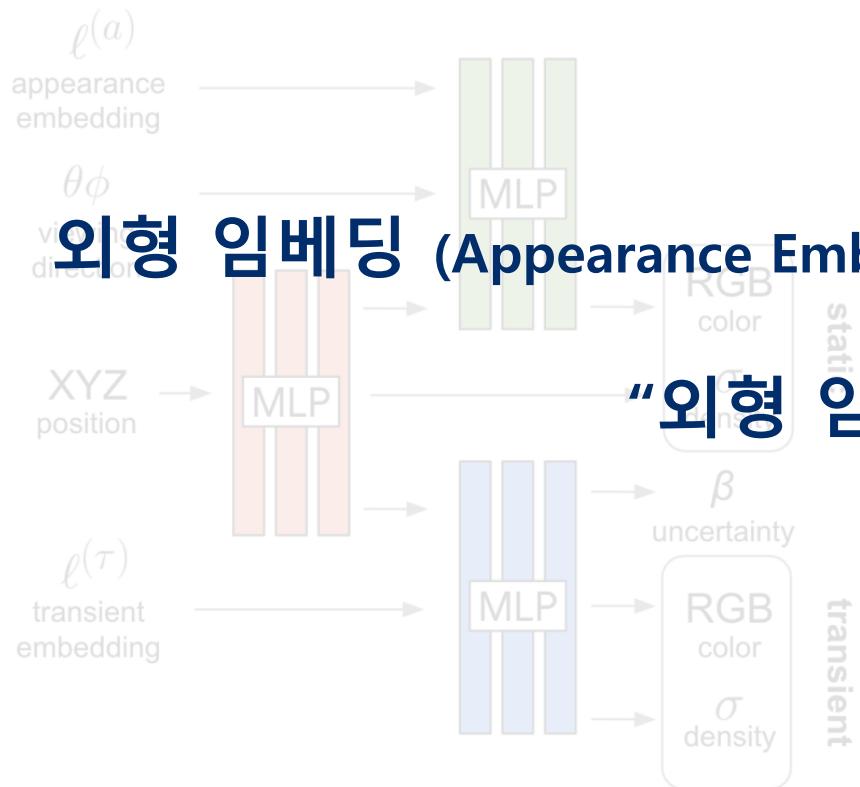
Related Work

NeRF in the Wild

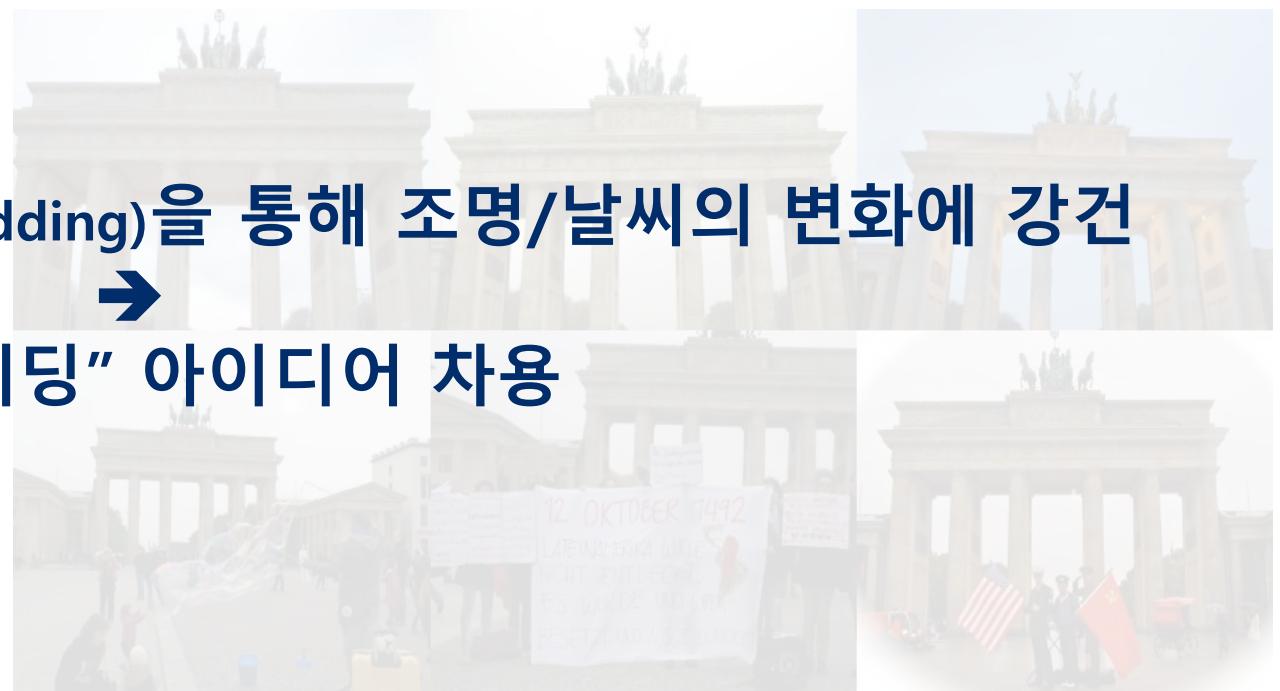


Related Work

NeRF in the Wild



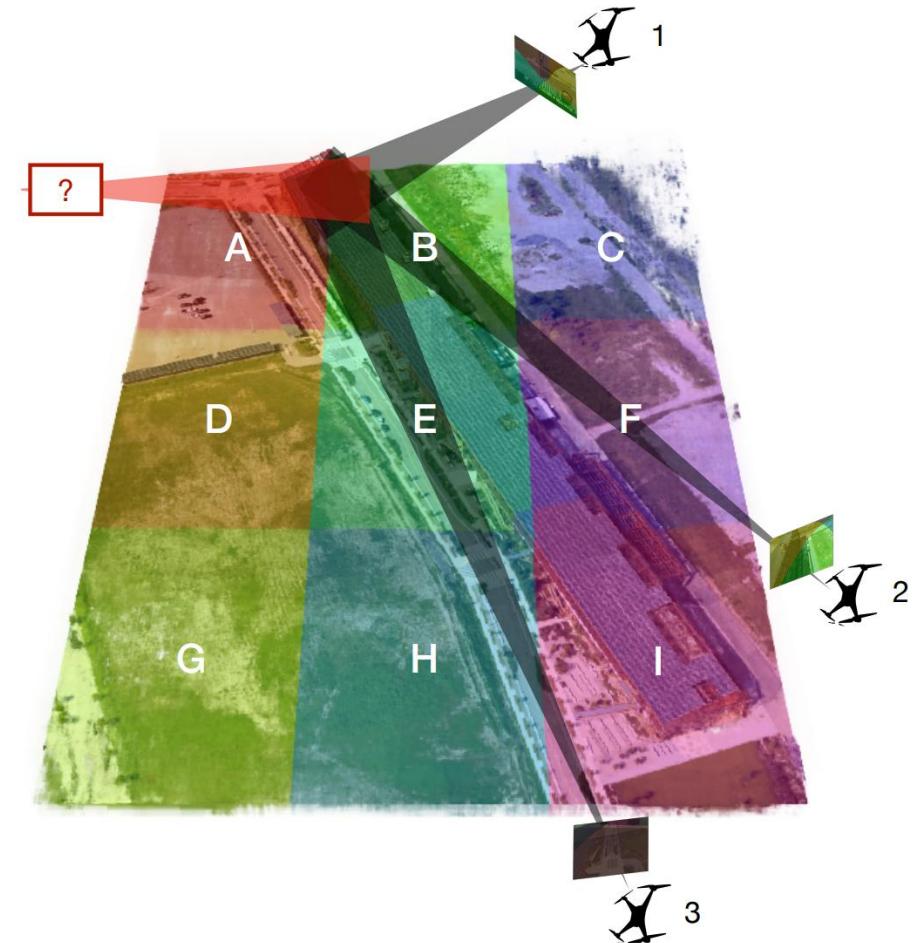
외형 임베딩 (Appearance Embedding)을 통해 조명/날씨의 변화에 강건



Method

Method - architecture

- 2D grid decomposition - efficiency
- Each cell has their own f^n (NeRF == submodule)
- Each cell's centroids $\mathbf{n} \in \mathcal{N} = (n_x, n_y, n_z)$
- **Input:** position, direction, appearance embedding vector $\mathbf{l}^{(a)}$
- **Output:** opacity, color



Method - architecture

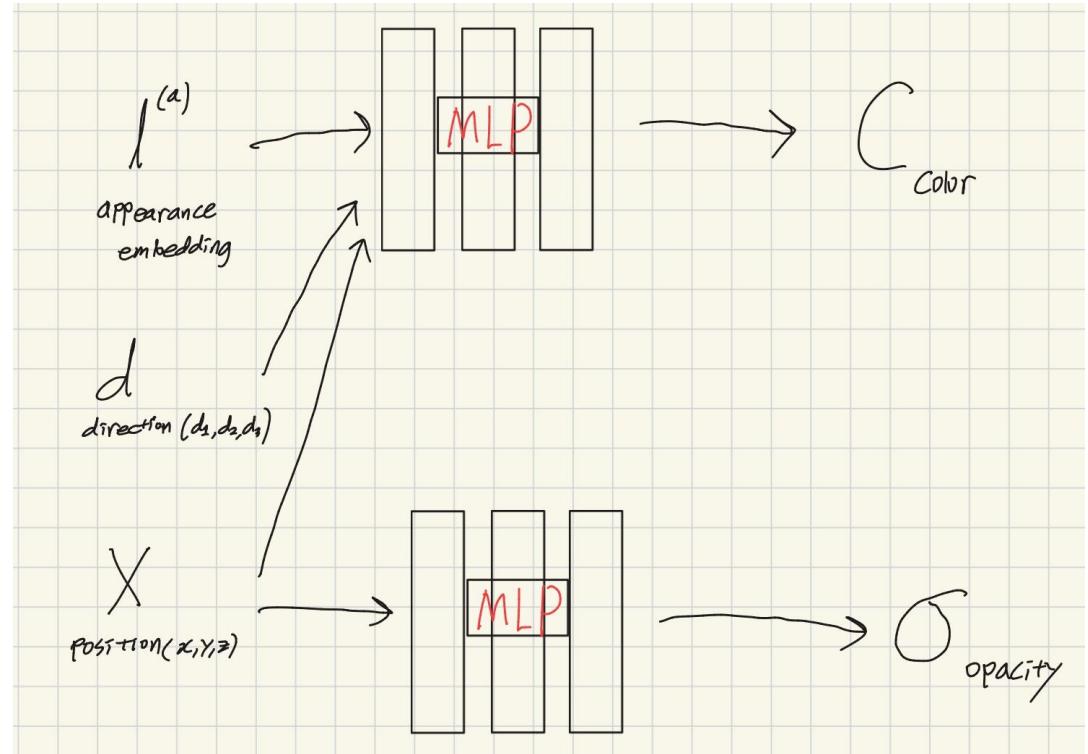
$$f^{\mathbf{n}}(\mathbf{x}) = \sigma \quad (1)$$

$$f^{\mathbf{n}}(\mathbf{x}, \mathbf{d}, l^{(a)}) = \mathbf{c} \quad (2)$$

$$\text{where } \mathbf{n} = \underset{n \in \mathcal{N}}{\operatorname{argmin}} \|n - \mathbf{x}\|^2 \quad (3)$$

$$\mathbf{n} \in \mathcal{N} = (n_x, n_y, n_z)$$

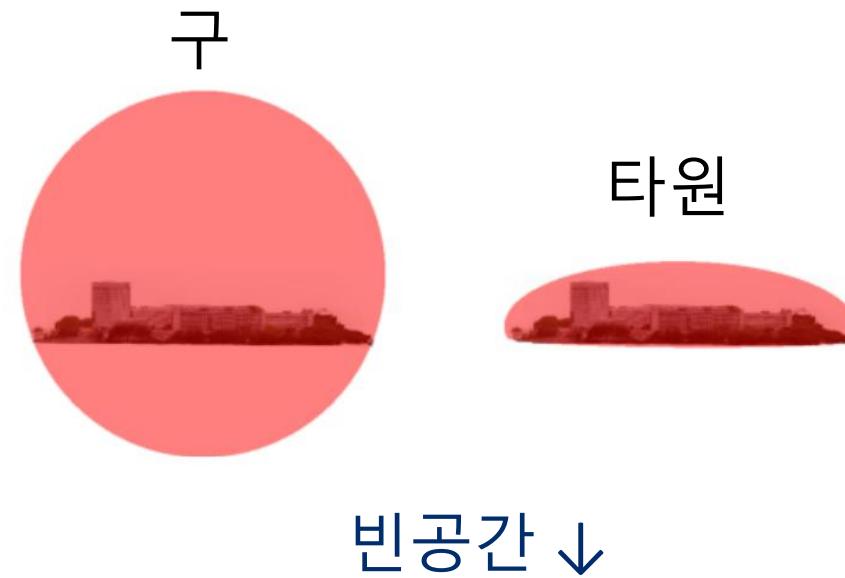
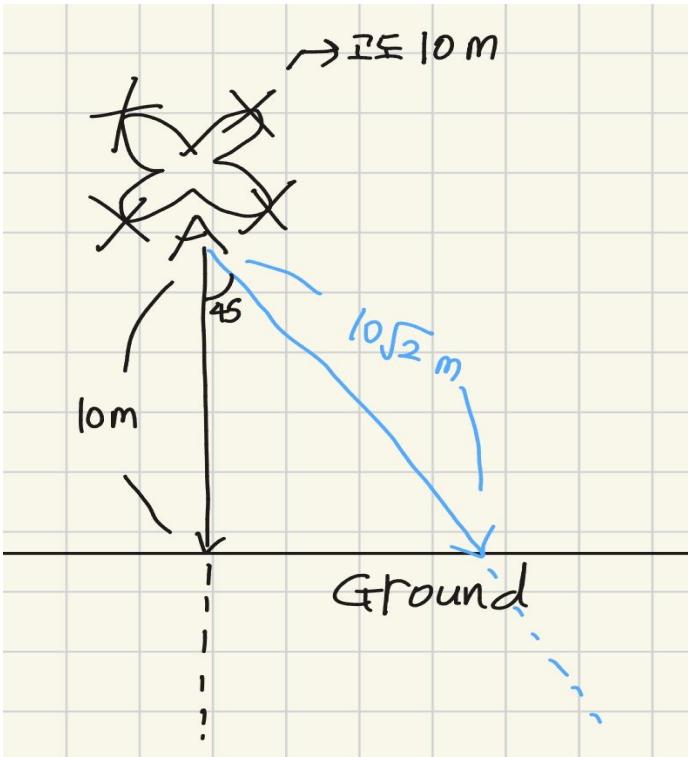
어떤 서브모듈을 사용할지 결정



Mega-NeRF Architecture

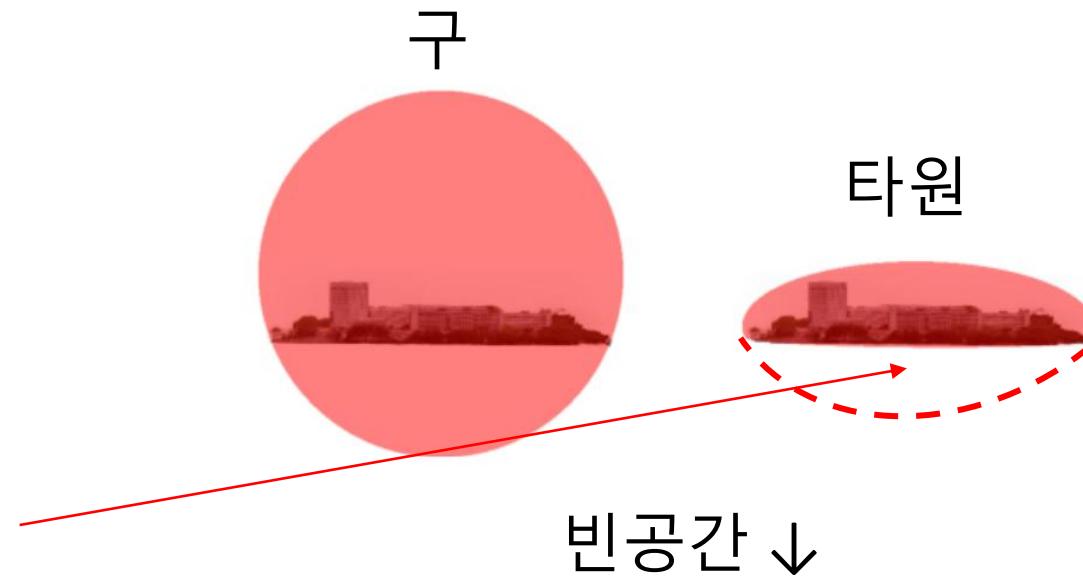
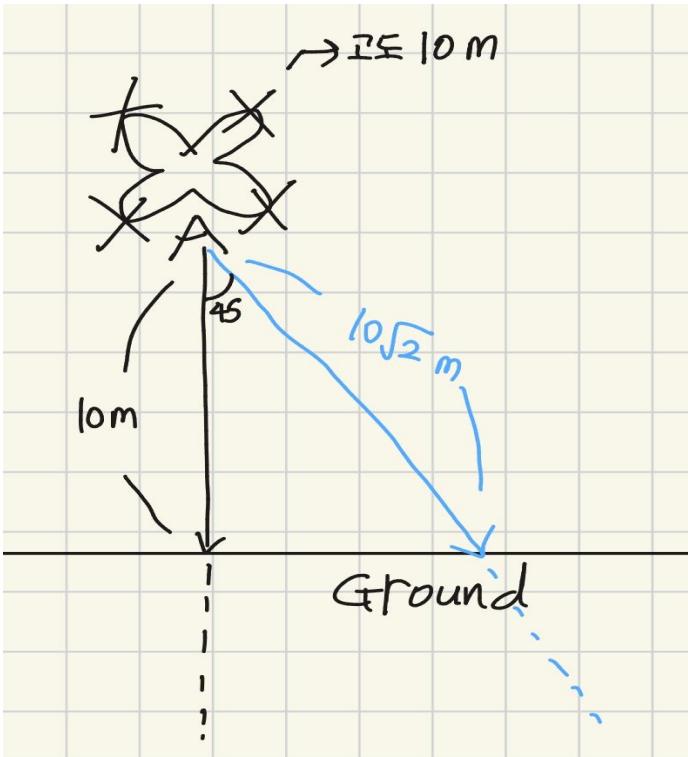
Method - architecture

- Terminate rays near **ground** level.
- Subdivide the scene into a **foreground** and a **background**



Method - architecture

- Terminate rays near **ground** level.
- Subdivide the scene into a **foreground** and a **background**

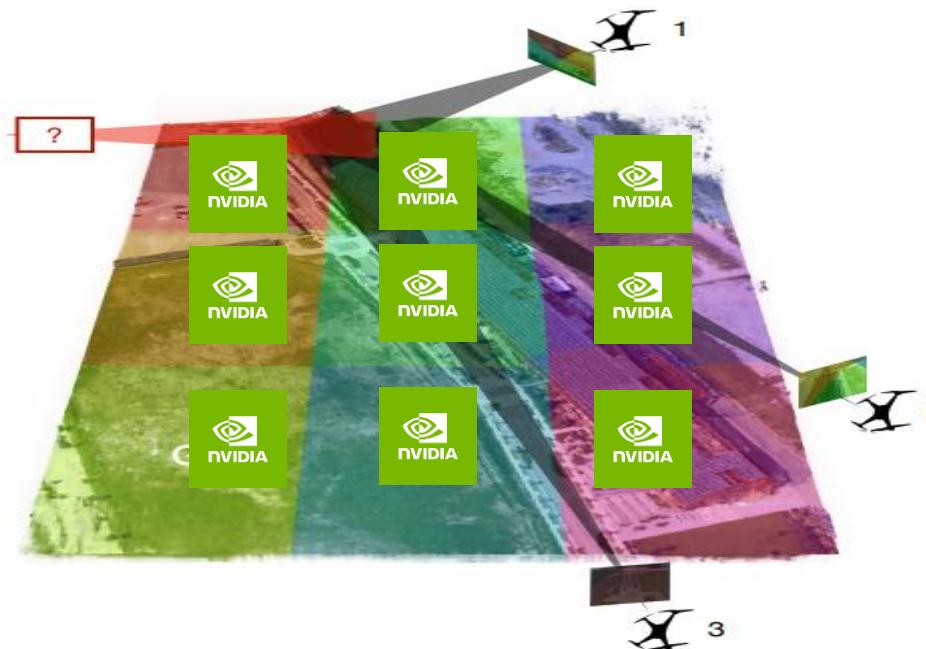


Method - Training

- Train each in parallel with **no inter-module communication**
- Add pixels to the trainset for only those spatial cells it intersects
- small **overlap** factor between cells (15% in our experiments)
- **Spatial Data Pruning**

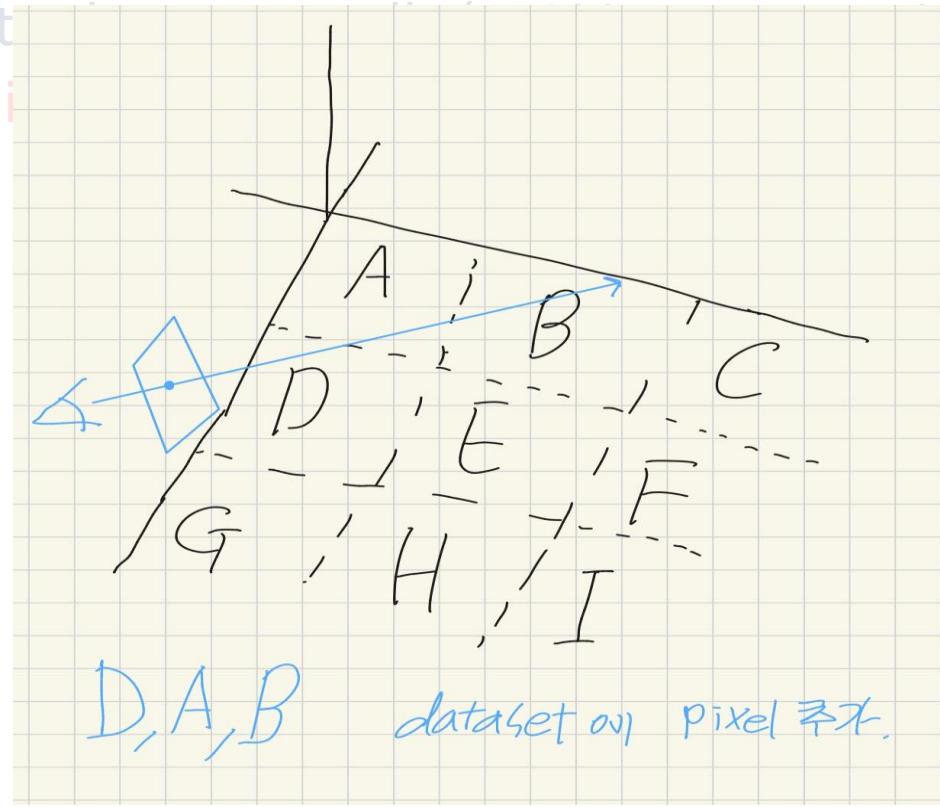
Method - Training

- Train each in parallel with **no inter-module communication**
- Add pixels to the trainset for only those spatial cells it intersects
- small **overlap** factor between cells (15% in our experiments)
- Spatial Data Pruning



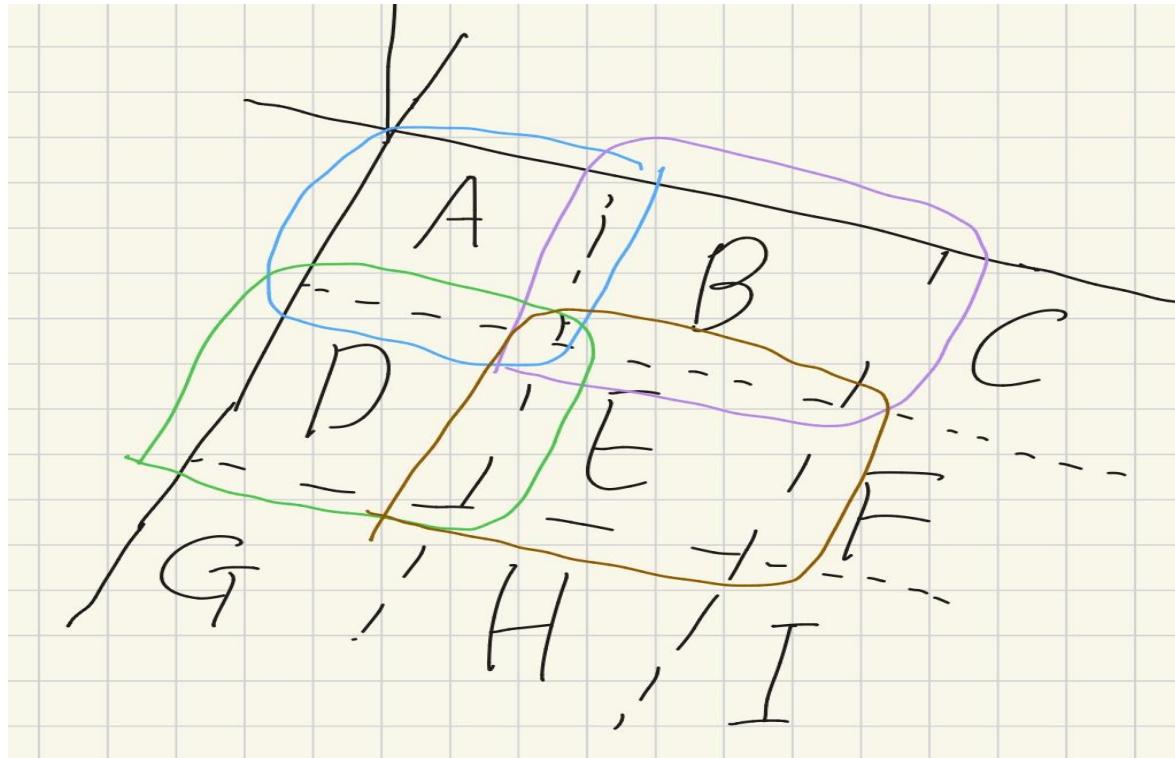
Method - Training

- Train each in parallel with no inter-module communication
- Add pixels to the trainset for only those spatial cells it intersects
- small overlap factor (e.g. 10%)
- Spatial Data Pruning



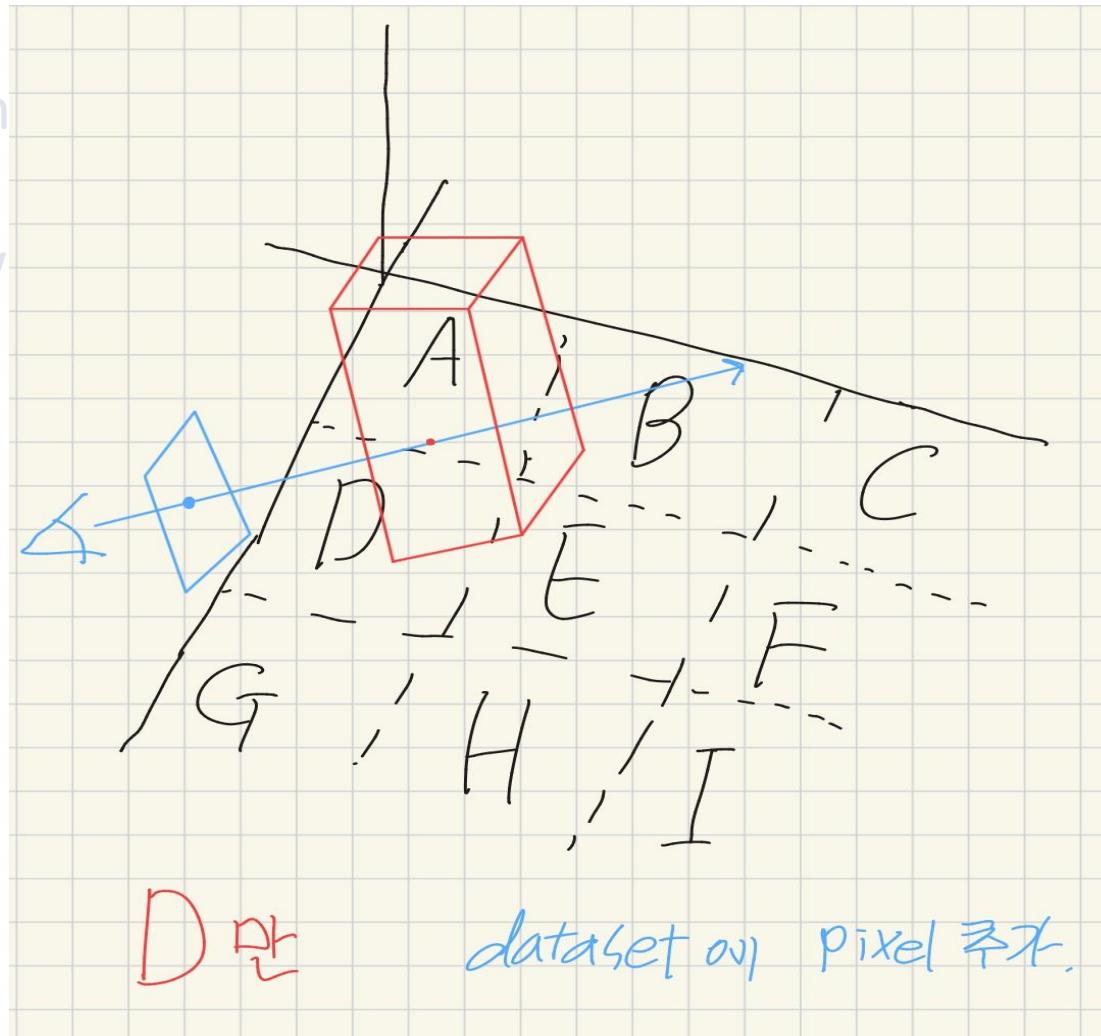
Method - Training

- Train each in parallel with **no inter-module communication**
- Add pixels to the trainset for only those spatial cells it intersects
- **small overlap factor between cells (15% in our experiments)**
- Spatial Data Pruning



Method - Training

- Train each in parallel with
- Add pixels to the trainset
- small **overlap** factor betw
- **Spatial Data Pruning**



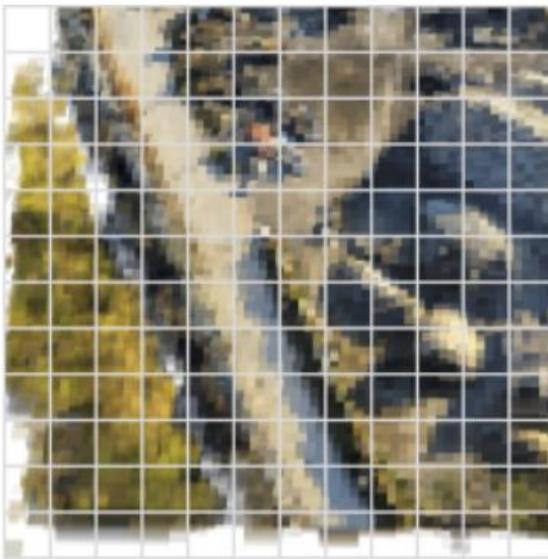
Method - Rendering

- Dynamic Octree
- Temporal coherence of interactive flythroughs
- Caching – LRU(Least Recently Used)
- Guided sampling

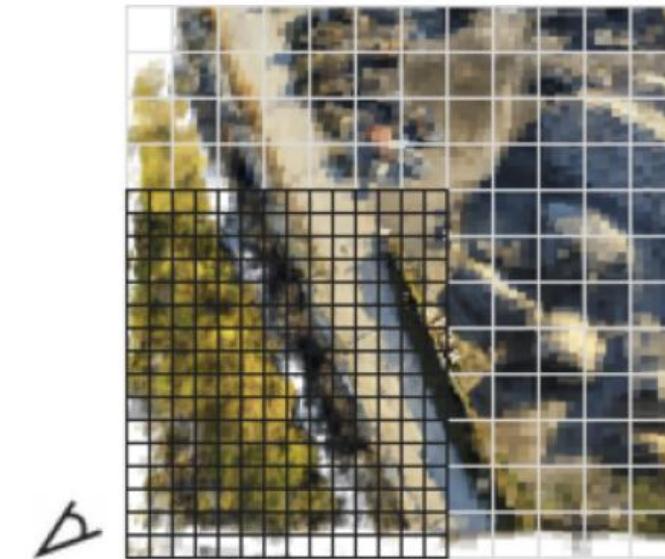
Method - Rendering

- Dynamic Octree
- Temporal coherence of interactive flythroughs
- Caching – LRU(Least Recently Used)

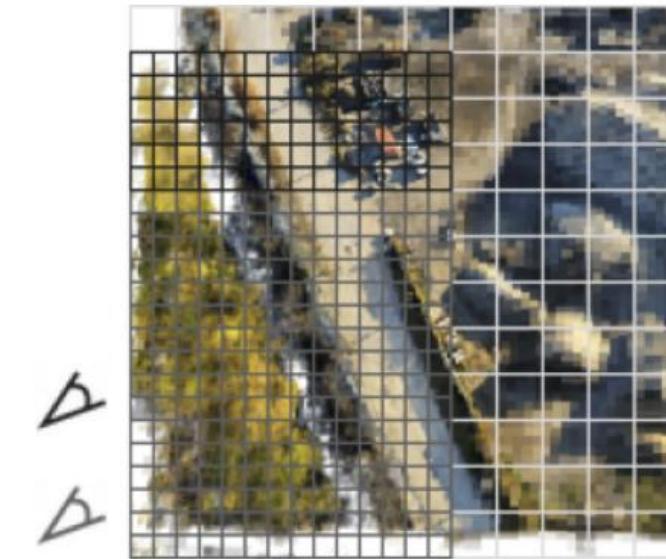
• Guided sampling



Sparse caching



View



Next View

Method - Rendering

- Dynamic Octree
- Temporal coherence of interactive flythroughs
- Caching – LRU(Least Recently Used)

Urban Scale environment

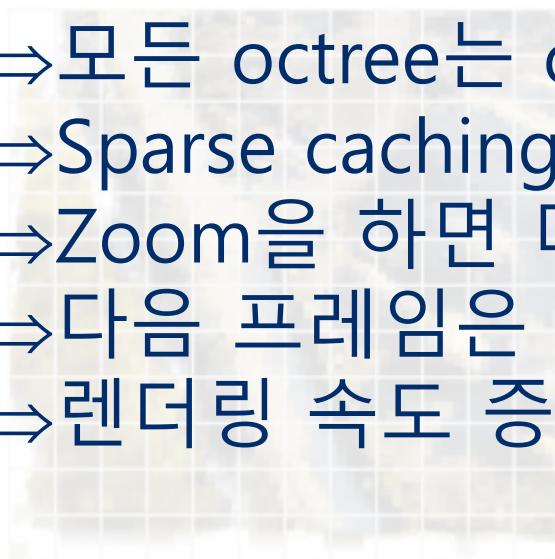
⇒ 모든 octree는 caching 불가

⇒ Sparse caching

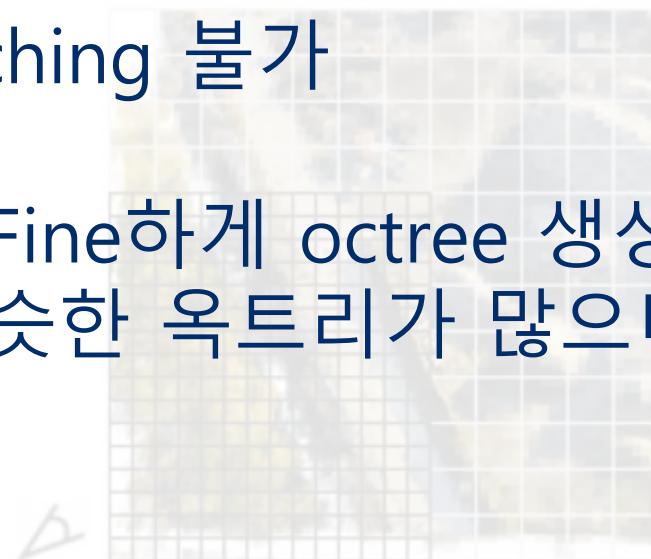
⇒ Zoom을 하면 더 Fine하게 octree 생성

⇒ 다음 프레임은 비슷한 옥트리가 많으니 연산량 감소

⇒ 렌더링 속도 증가



Sparse caching



View



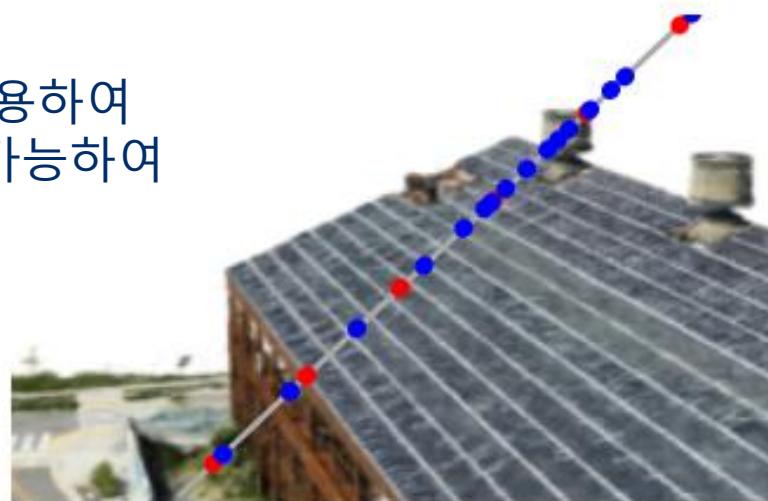
Next View

Method - Rendering

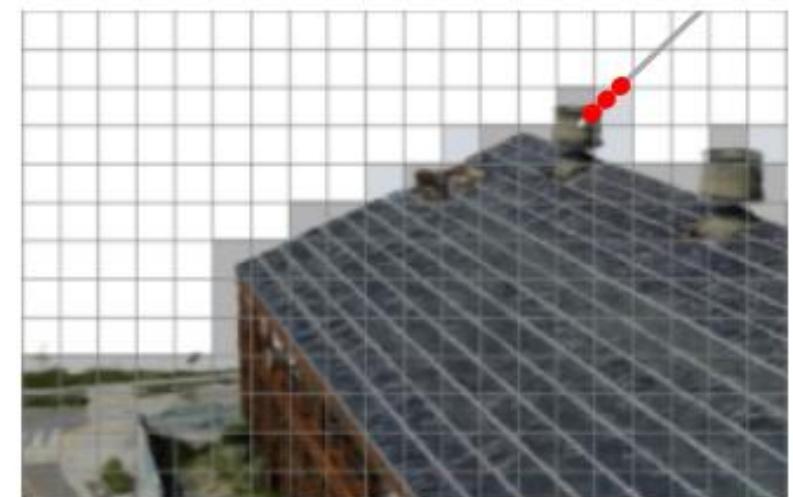
- Temporal coherence of interactive flythroughs
- Caching – LRU(Least Recently Used)
- Guided sampling

*Octree: 3D 공간을 8개의 작은 정육면체로 재귀적으로 계속 쪼개서 데이터를 관리하는 트리(Tree) 자료구조 (RGB와 Opacity 값을 보관)

Octree의 opacity 값을 활용하여
octree가 비었는지 확인 가능하여
샘플링을 줄일 수 있다.



Standard Hierarchical Sampling
NeRF

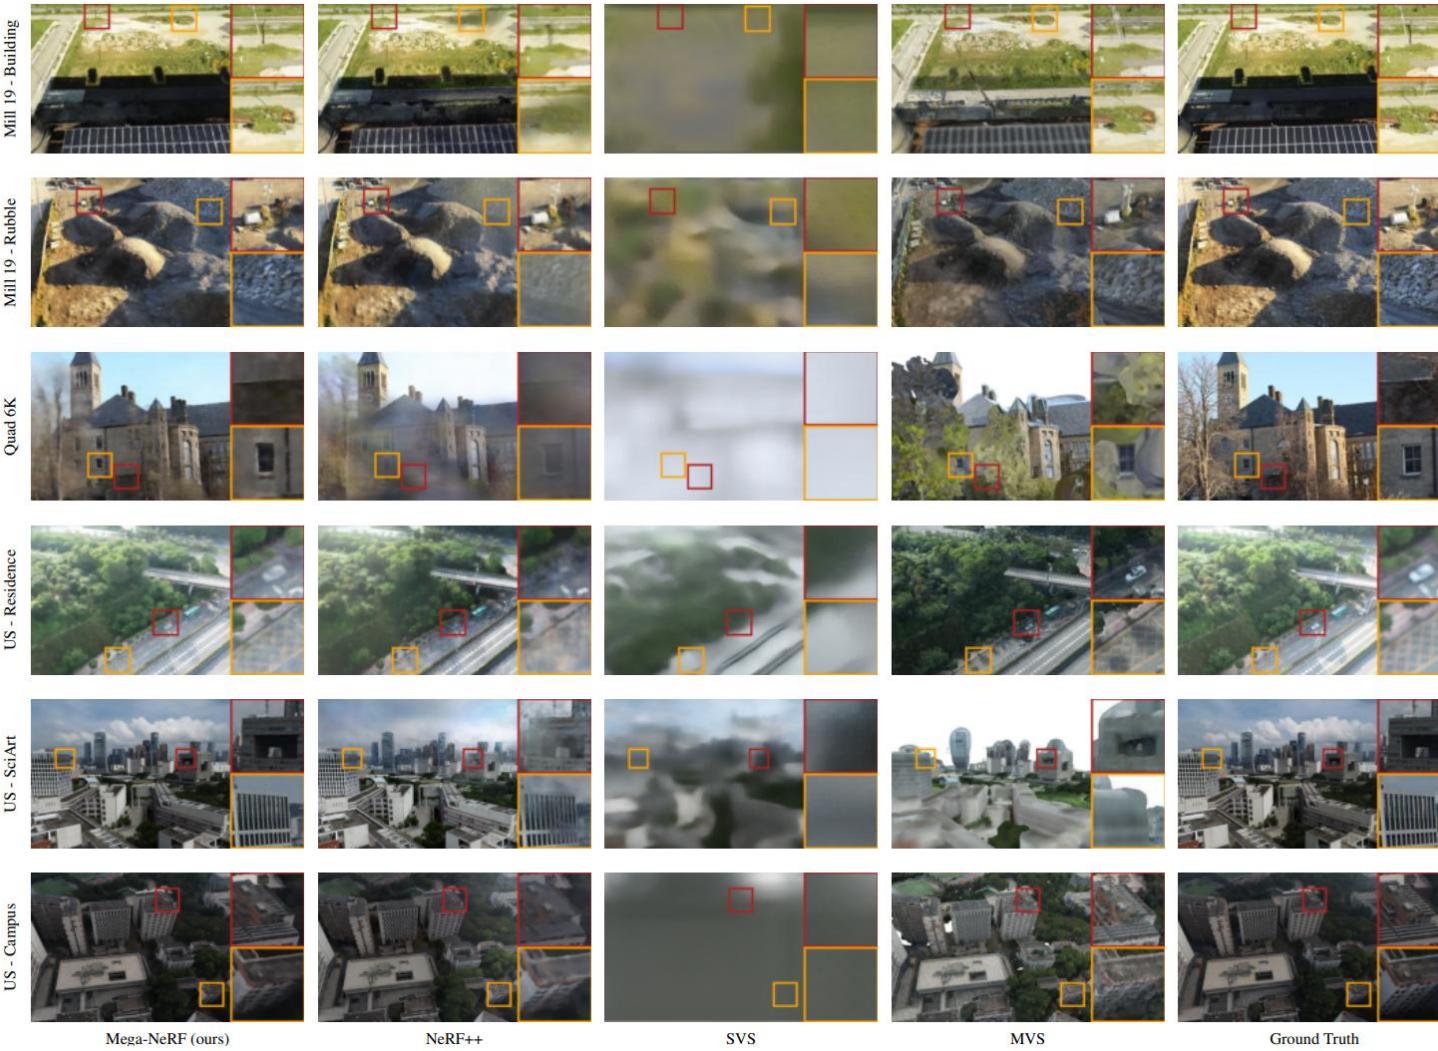


Guided Sampling
Mega-NeRF

Experiment

Experiment

Training



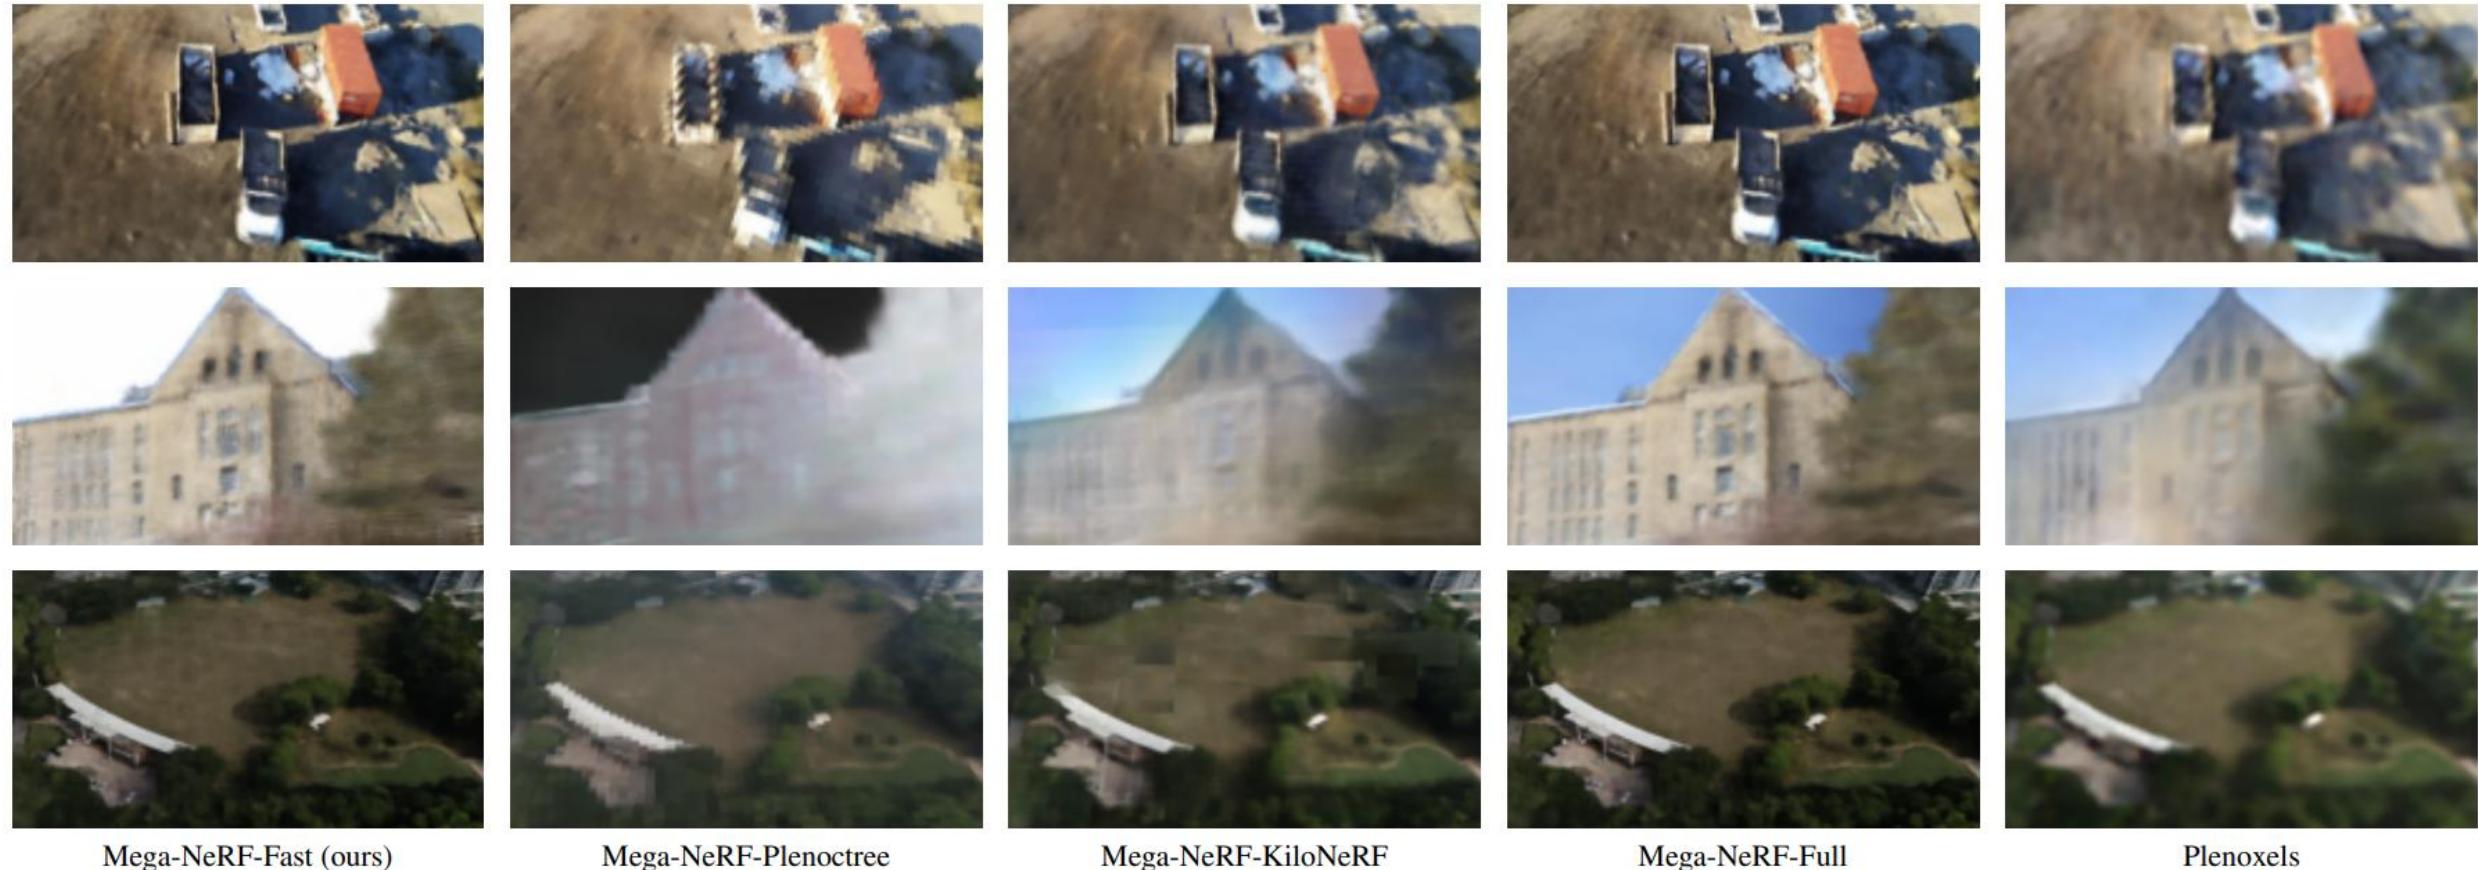
Experiment

Training

	Mill 19 - Building				Mill 19 - Rubble				Quad 6k			
	↑PSNR	↑SSIM	↓LPIPS	↓Time (h)	↑PSNR	↑SSIM	↓LPIPS	↓Time(h)	↑PSNR	↑SSIM	↓LPIPS	↓Time(h)
NeRF	19.54	0.525	0.512	59:51	21.14	0.522	0.546	60:21	16.75	0.559	0.616	62:48
NeRF++	19.48	0.520	0.514	89:02	20.90	0.519	0.548	90:42	16.73	0.560	0.611	90:34
SVS	12.59	0.299	0.778	38:17	13.97	0.323	0.788	37:33	11.45	0.504	0.637	29:48
DeepView	13.28	0.295	0.751	31:20	14.47	0.310	0.734	32:11	11.34	0.471	0.708	19:51
MVS	16.45	0.451	0.545	32:29	18.59	0.478	0.532	31:42	11.81	0.425	0.594	18:55
Mega-NeRF	20.93	0.547	0.504	29:49	24.06	0.553	0.516	30:48	18.13	0.568	0.602	39:43
	UrbanScene3D - Residence				UrbanScene3D - Sci-Art				UrbanScene3D - Campus			
	↑PSNR	↑SSIM	↓LPIPS	↓Time (h)	↑PSNR	↑SSIM	↓LPIPS	↓Time(h)	↑PSNR	↑SSIM	↓LPIPS	↓Time(h)
NeRF	19.01	0.593	0.488	62:40	20.70	0.727	0.418	60:15	21.83	0.521	0.630	61:56
NeRF++	18.99	0.586	0.493	90:48	20.83	0.755	0.393	95:00	21.81	0.520	0.630	93:50
SVS	16.55	0.388	0.704	77:15	15.05	0.493	0.716	59:58	13.45	0.356	0.773	105:01
DeepView	13.07	0.313	0.767	30:30	12.22	0.454	0.831	31:29	13.77	0.351	0.764	33:08
MVS	17.18	0.532	0.429	69:07	14.38	0.499	0.672	73:24	16.51	0.382	0.581	96:01
Mega-NeRF	22.08	0.628	0.489	27:20	25.60	0.770	0.390	27:39	23.42	0.537	0.618	29:03

Experiment

Rendering



Experiment

Rendering

best second-best	Mill 19					Quad 6k					UrbanScene3D				
	↑PSNR	↑SSIM	↓LPIPS	Preprocess	Render	↑PSNR	↑SSIM	↓LPIPS	Preprocess	Render	↑PSNR	↑SSIM	↓LPIPS	Preprocess	Render
			Time (h)	Time (s)				Time (h)	Time (s)				Time (h)	Time (s)	
Mega-NeRF-Plenoctree	16.27	0.430	0.621	1:26	0.031	13.88	0.589	0.427	1:33	0.010	16.41	0.498	0.530	1:07	0.025
Mega-NeRF-KiloNeRF	21.85	0.521	0.512	30:03	0.784	20.61	0.652	0.356	27:33	1.021	21.11	0.542	0.453	34:00	0.824
Mega-NeRF-Full	22.96	0.588	0.452	-	101	21.52	0.676	<u>0.355</u>	-	174	24.92	0.710	0.393	-	122
Plenoxels	19.32	0.476	0.592	-	0.482	18.61	0.645	0.411	-	<u>0.194</u>	20.06	0.608	0.503	-	0.531
Mega-NeRF-Initial	17.41	0.447	0.570	1:08	<u>0.235</u>	14.30	0.585	0.386	1:31	0.214	17.22	0.527	0.506	1:10	<u>0.221</u>
Mega-NeRF-Dynamic	<u>22.34</u>	<u>0.573</u>	<u>0.464</u>	1:08	3.96	<u>20.84</u>	<u>0.658</u>	0.342	1:31	2.91	<u>23.99</u>	<u>0.691</u>	<u>0.408</u>	1:10	3.219

Experiment

Ablation Study

	Mill 19			Quad 6k			UrbanScene3D		
	↑PSNR	↑SSIM	↓LPIPS	↑PSNR	↑SSIM	↓LPIPS	↑PSNR	↑SSIM	↓LPIPS
Mega-NeRF-no-embed	20.42	0.500	0.561	16.16	0.544	0.643	19.45	0.587	0.545
Mega-NeRF-embed-only	21.48	0.494	0.566	17.91	0.559	0.638	22.79	0.611	0.537
Mega-NeRF-no-bounds	22.14	0.534	0.522	18.02	0.565	0.616	23.42	0.636	0.511
Mega-NeRF-dense	21.63	0.504	0.551	17.94	0.562	0.627	22.44	0.605	0.558
Mega-NeRF-joint	21.10	0.490	0.574	17.43	0.560	0.616	21.45	0.595	0.567
Mega-NeRF	22.34	0.540	0.518	18.08	0.566	0.602	23.60	0.641	0.504

Contribution

Contribution

- Presents a novel rendering method that exploits **temporal coherence**.
- Presents a new **large-scale dataset**.
- Divide the model into multiple **submodules** and **train** in a **fully parallelizable** manner.